

## 76. Ethical Frameworks for Strategic AI

Consequentialism, deontology, and virtue ethics applied to AI decisions, with formal social welfare functions for resource allocation.



# 77. Ethical Frameworks for Strategic AI

## Learning objectives

- Distinguish consequentialist, deontological, and virtue-ethics perspectives on strategic AI decisions.
- Formalize utilitarian, Rawlsian, and Nash social welfare functions and compute their optima.
- Implement and compare welfare criteria for multi-agent resource allocation in R.
- Identify Pareto-efficient allocations and locate welfare-maximizing points on the Pareto frontier.

## 77.1. Motivation

An autonomous system must allocate a scarce resource — say, bandwidth, computing capacity, or medical supplies — among multiple agents. A purely strategic analysis tells us what equilibrium outcomes look like, but it says nothing about which equilibrium is *desirable*. Should we maximize total benefit (utilitarianism)? Protect the worst-off agent (Rawlsian maximin)? Find a balanced compromise (Nash bargaining)?

These questions are not merely philosophical. Every multi-agent system that resolves conflicts among stakeholders implicitly embeds an ethical framework in its objective function. Making the framework explicit allows designers to reason about trade-offs, justify choices, and audit outcomes.

## 77.2. Theory

### 77.2.1. Three ethical lenses

**Consequentialism** evaluates actions by their outcomes. In a game-theoretic setting, the most common consequentialist criterion is *utilitarianism*: choose the outcome that maximizes the sum of payoffs across all agents.

**Deontology** evaluates actions by whether they conform to rules or duties, regardless of consequences. In mechanism design this corresponds to *constraints* — for example, requiring that no agent is made worse off than their outside option (individual rationality) or that the mechanism treats agents symmetrically.

**Virtue ethics** asks what a well-functioning agent *would* do. In AI alignment, this maps to designing agents that exhibit desirable dispositions — honesty, fairness, cooperativeness — as stable character traits rather than case-by-case calculations.

### 77.2.2. Social welfare functions

Given  $n$  agents with utilities  $u_1, u_2, \dots, u_n$  over an allocation  $x$ , we define three standard social welfare functions (SWFs):

**Utilitarian SWF** — maximize total welfare:

$$W_U(x) = \sum_{i=1}^n u_i(x) \quad (77.1)$$

**Rawlsian (maximin) SWF** — maximize the welfare of the worst-off agent:

$$W_R(x) = \min_i u_i(x) \quad (77.2)$$

**Nash SWF** — maximize the product of utilities (equivalent to Nash bargaining with equal bargaining power):

$$W_N(x) = \prod_{i=1}^n u_i(x) \quad (77.3)$$

#### **i** Pareto efficiency

An allocation  $x$  is **Pareto efficient** if there is no alternative  $x'$  such that  $u_i(x') \geq u_i(x)$  for all  $i$  with strict inequality for at least one. All three SWF optima lie on the Pareto frontier, but they generally select different points.

### 77.2.3. Impossibility and trade-offs

Arrow's impossibility theorem and its descendants show that no social welfare function can satisfy all desirable axioms simultaneously. In practice, the choice of SWF reflects a value judgement about how to weigh efficiency against equity. Utilitarianism favours efficiency, the Rawlsian criterion favours equity, and the Nash criterion offers a middle ground that is scale-invariant and satisfies the axioms of Nash bargaining.

## 77.3. Implementation in R

We consider a resource allocation problem: divide a budget  $B = 100$  among three agents whose utilities are concave functions of their share.

```
# Agent utility functions (concave - diminishing returns)
utility <- function(x, alpha) {
  # CES-style utility: x^alpha, alpha in (0, 1) gives concavity
  x^alpha
}

# Agent parameters: different elasticities
alphas <- c(0.3, 0.5, 0.8)
budget <- 100
n_agents <- 3
```

```
# Social welfare functions
swf_utilitarian <- function(alloc, alphas) {
  sum(mapply(utility, alloc, alphas))
}

swf_rawlsian <- function(alloc, alphas) {
  min(mapply(utility, alloc, alphas))
}

swf_nash <- function(alloc, alphas) {
  utils <- mapply(utility, alloc, alphas)
  if (any(utils <= 0)) return(0)
  prod(utils)
}
```

### 77.3.1. Grid search over allocations

```
# Generate feasible allocations on a grid ( $x_1 + x_2 + x_3 = \text{budget}$ )
step <- 2
grid <- expand.grid(
  x1 = seq(step, budget - 2 * step, by = step),
  x2 = seq(step, budget - 2 * step, by = step)
) |>
  mutate(x3 = budget - x1 - x2) |>
  filter(x3 >= step)

# Compute utilities and welfare for each allocation
results <- grid |>
  rowwise() |>
  mutate(
    u1 = utility(x1, alphas[1]),
    u2 = utility(x2, alphas[2]),
    u3 = utility(x3, alphas[3]),
    W_util = u1 + u2 + u3,
    W_rawls = min(u1, u2, u3),
    W_nash = u1 * u2 * u3
  ) |>
  ungroup()

# Find optima
opt_util <- results |> slice_max(W_util, n = 1, with_ties = FALSE)
opt_rawls <- results |> slice_max(W_rawls, n = 1, with_ties = FALSE)
opt_nash <- results |> slice_max(W_nash, n = 1, with_ties = FALSE)

cat("Utilitarian optimum:  x =", c(opt_util$x1, opt_util$x2, opt_util$x3), "\n")
```

```
#> Utilitarian optimum:  x = 2 2 96
```

```
cat("Rawlsian optimum:    x =", c(opt_rawls$x1, opt_rawls$x2, opt_rawls$x3), "\n")
```

```
#> Rawlsian optimum:    x = 80 14 6
```

```
cat("Nash optimum:      x =", c(opt_nash$x1, opt_nash$x2, opt_nash$x3), "\n")
```

```
#> Nash optimum:      x = 18 32 50
```

### 77.3.2. Pareto frontier with welfare optima

```
# Identify Pareto-efficient allocations (no other allocation dominates)
is_dominated <- function(i, df) {
  u <- c(df$u1[i], df$u2[i], df$u3[i])
  any(
    df$u1 >= u[1] & df$u2 >= u[2] & df$u3 >= u[3] &
    (df$u1 > u[1] | df$u2 > u[2] | df$u3 > u[3])
  )
}

pareto_mask <- !apply(seq_len(nrow(results)), is_dominated, df = results)
pareto <- results[pareto_mask, ]
```

```
optima_points <- bind_rows(
  opt_util |> mutate(criterion = "Utilitarian"),
  opt_rawls |> mutate(criterion = "Rawlsian"),
  opt_nash |> mutate(criterion = "Nash")
)

p1 <- ggplot() +
  geom_point(data = results, aes(x = u1, y = u2),
    colour = "grey80", size = 0.5, alpha = 0.3) +
  geom_point(data = pareto, aes(x = u1, y = u2),
    colour = "grey40", size = 1, alpha = 0.6) +
  geom_point(data = optima_points, aes(x = u1, y = u2, colour = criterion),
    size = 4, shape = 17) +
  scale_colour_manual(
    values = c("Utilitarian" = okabe_ito[1],
      "Rawlsian" = okabe_ito[2],
      "Nash" = okabe_ito[3]),
    name = "Welfare criterion"
  ) +
  labs(
    title = "Pareto Frontier with Social Welfare Optima",
    x = expression(u[1] ~ "(Agent 1 utility)"),
    y = expression(u[2] ~ "(Agent 2 utility)")
  ) +
  theme_publication()
```

```
p1
save_pub_fig(p1, "ethics-pareto-frontier", width = 7, height = 5)
```




Figure 77.1: Pareto frontier projected onto agents 1 and 2 utility space. The utilitarian (orange), Rawlsian (blue), and Nash (green) optima select different points on the frontier, reflecting different equity-efficiency trade-offs.

### 77.3.3. Welfare comparison across methods

```
comparison <- optima_points |>
  select(criterion, u1, u2, u3) |>
  pivot_longer(cols = c(u1, u2, u3),
               names_to = "agent", values_to = "utility") |>
  mutate(agent = recode(agent, u1 = "Agent 1", u2 = "Agent 2", u3 = "Agent 3"))

p2 <- ggplot(comparison, aes(x = criterion, y = utility, fill = agent)) +
  geom_col(position = position_dodge(width = 0.7), width = 0.6) +
  scale_fill_manual(
    values = c("Agent 1" = okabe_ito[1],
              "Agent 2" = okabe_ito[2],
              "Agent 3" = okabe_ito[3]),
    name = "Agent"
  ) +
  labs(
    title = "Welfare Comparison Across Allocation Methods",
    x = "Welfare criterion",
    y = "Individual utility"
  ) +
  theme_publication()
```

```
p2
save_pub_fig(p2, "ethics-welfare-comparison", width = 7, height = 5)
```




Figure 77.2.: Individual agent utilities under each welfare criterion. The utilitarian optimum concentrates resources on the agent with highest marginal utility (agent 3), while the Rawlsian criterion equalizes utilities across agents.

## 77.4. Worked example

Consider three agents sharing a budget of  $B = 100$  with utility functions  $u_i(x_i) = x_i^{\alpha_i}$  where  $\alpha_1 = 0.3$ ,  $\alpha_2 = 0.5$ , and  $\alpha_3 = 0.8$ . Agent 3 has the highest elasticity — each additional unit of resource generates more utility for agent 3 than for the others.

**Step 1 — Utilitarian optimum.** The utilitarian criterion maximizes  $\sum u_i$ . Because agent 3's marginal utility decreases most slowly ( $\alpha_3 = 0.8$ ), the utilitarian solution allocates the most resources to agent 3.

```
cat("=== Worked Example: 3-Agent Resource Allocation ===\n\n")
```

```
#> === Worked Example: 3-Agent Resource Allocation ===
```

```
cat("Agent parameters (alpha):", alphas, "\n")
```

```
#> Agent parameters (alpha): 0.3 0.5 0.8
```

```
cat("Budget:", budget, "\n\n")
```

```
#> Budget: 100
```



```
# Report allocations and utilities for each criterion
for (lbl in c("Utilitarian", "Rawlsian", "Nash")) {
  row <- optima_points |> filter(criterion == lbl)
  cat(sprintf("%s optimum:\n", lbl))
  cat(sprintf("  Allocation: (%.0f, %.0f, %.0f)\n", row$x1, row$x2, row$x3))
  cat(sprintf("  Utilities:   (%.3f, %.3f, %.3f)\n", row$u1, row$u2, row$u3))
  cat(sprintf("  Sum: %.3f | Min: %.3f | Product: %.3f\n\n",
              row$u1 + row$u2 + row$u3,
              min(row$u1, row$u2, row$u3),
              row$u1 * row$u2 * row$u3))
}
```

```
#> Utilitarian optimum:
#>   Allocation: (2, 2, 96)
#>   Utilities:  (1.231, 1.414, 38.532)
#>   Sum: 41.177 | Min: 1.231 | Product: 67.087
#>
#> Rawlsian optimum:
#>   Allocation: (80, 14, 6)
#>   Utilities:  (3.723, 3.742, 4.193)
#>   Sum: 11.658 | Min: 3.723 | Product: 58.413
#>
#> Nash optimum:
#>   Allocation: (18, 32, 50)
#>   Utilities:  (2.380, 5.657, 22.865)
#>   Sum: 30.902 | Min: 2.380 | Product: 307.845
```

**Step 2 — Rawlsian optimum.** The maximin criterion equalizes utilities as much as possible. Because agent 1 has the most concave utility function, it requires a larger share of the budget to achieve the same utility as the others.

**Step 3 — Nash optimum.** The Nash product criterion finds a middle ground. It is the unique allocation satisfying the axioms of symmetry, Pareto efficiency, independence of irrelevant alternatives, and scale invariance from Nash’s bargaining theory.

**Interpretation.** The three criteria represent genuine ethical disagreements. A designer choosing among them is making a normative choice: how much inequality is acceptable in exchange for higher total welfare? There is no objectively correct answer, but formalizing the alternatives allows transparent discussion and principled selection.

## 77.5. Extensions

- **Weighted welfare functions.** Assign different weights  $w_i$  to agents:  $W_U(x) = \sum w_i u_i(x)$ . This can represent priority given to disadvantaged groups or domain-specific considerations.
- **Mechanism design** (Chapter 67) connects welfare functions to incentive-compatible mechanisms: can we design rules that achieve a desired welfare criterion even when agents act strategically?
- **Fairness in ML** (Chapter 79) applies similar trade-offs to machine learning classifiers, where “agents” are demographic groups and “utility” is prediction accuracy.

- For philosophical foundations, see Sen (1999) and Rawls (1971). For the game-theoretic connection, Moulin (2004) provides an axiomatic treatment of fair division.

## Exercises

1. **Weighted utilitarianism.** Modify the grid search to maximize  $W(x) = 2u_1(x) + u_2(x) + u_3(x)$ , giving agent 1 double weight. How does the optimal allocation change compared to equal-weight utilitarianism? Plot the result on the Pareto frontier.
2. **Four agents.** Extend the implementation to four agents with  $\alpha = (0.2, 0.4, 0.6, 0.9)$  and budget  $B = 200$ . Compare the utilitarian and Rawlsian allocations. Which agent gains the most from switching to the Rawlsian criterion?
3. **Leximin refinement.** The Rawlsian criterion only considers the worst-off agent. Implement the *leximin* refinement, which first maximizes the minimum utility, then (among allocations achieving this maximum) maximizes the second-lowest utility, and so on. Does leximin ever differ from plain maximin in the three-agent example?

Solutions appear in Appendix H.