

I. Glossary

Key terms and definitions used throughout this book.

J. Glossary

- Action** A choice available to a player at a decision point. In a simultaneous game, the terms “action” and “strategy” are often used interchangeably.
- Agent-based model** A computational model in which autonomous agents interact according to specified rules, producing emergent macro-level behaviour. See Chapter 35.
- Backward induction** A solution method for extensive-form games of perfect information. Starting from terminal nodes and working backwards, each player’s optimal action is determined at every decision node. The resulting strategy profile is a subgame perfect equilibrium. See Chapter 13.
- Battle of the Sexes** A 2x2 coordination game with two pure-strategy Nash equilibria and one mixed equilibrium. The players prefer to coordinate but disagree on which outcome is better.
- Bayesian game** A game in which players have private information (types) drawn from a known prior distribution. Players maximise expected payoffs given their beliefs about opponents’ types. See Chapter 15.
- Bayesian Nash equilibrium (BNE)** A strategy profile in a Bayesian game where each type of each player maximises expected payoff given the prior distribution over opponents’ types and their strategies.
- Best response** A strategy s_i^* that maximises player i ’s payoff given the strategies of all other players. A Nash equilibrium is a profile where every player’s strategy is a best response.
- Characteristic function** In cooperative game theory, a function $v : 2^N \rightarrow \mathbb{R}$ that assigns a worth to every coalition $S \subseteq N$. See Chapter 19.
- Chicken (Hawk-Dove)** A 2x2 anti-coordination game modelling brinkmanship. Each player prefers the other to yield. It has two asymmetric pure-strategy NE and one mixed NE.
- Common knowledge** A fact is common knowledge if all players know it, all players know that all players know it, and so on ad infinitum. The structure of the game is typically assumed to be common knowledge.
- Cooperative game theory** The branch of game theory that studies what coalitions of players can achieve and how the joint gains should be allocated. Focuses on outcomes rather than strategies. See Chapter 19.
- Core** The set of feasible payoff allocations in a cooperative game that no coalition can improve upon. An allocation is in the core if no subset of players can do better by breaking away.
- Discount factor** A parameter $\delta \in [0, 1)$ representing the weight a player places on future payoffs relative to current payoffs. Central to repeated game analysis. See Chapter 17.
- Dominant strategy** A strategy that yields a payoff at least as high as any alternative regardless of opponents’ actions (weak dominance) or strictly higher (strict dominance).
- Evolutionary stable strategy (ESS)** A strategy that, once prevalent in a population, cannot be invaded by any rare mutant strategy. Formally, a strategy s^* is an ESS if it is a best response to itself and, against any alternative best response, does strictly better against the mutant than the mutant does against itself. See [?@sec-evolutionary-stable-strategies](#).
- Extensive form** A representation of a game as a tree, specifying the order of moves, information available at each decision point, and payoffs at terminal nodes. See Chapter 13.

- Feasible payoff set** The convex hull of all achievable payoff vectors in a game. In repeated games, the folk theorem concerns the subset of feasible payoffs that are individually rational.
- Folk theorem** A family of results stating that in infinitely repeated games with sufficiently patient players (δ close to 1), any feasible and individually rational payoff vector can be sustained as a Nash equilibrium outcome. See Chapter 17.
- Game tree** The tree representation of an extensive-form game, with nodes representing decision points and edges representing actions.
- Grim Trigger** A strategy in repeated games that cooperates until the opponent defects, then defects forever. The simplest strategy that can sustain cooperation via the threat of permanent punishment.
- Imperfect information** A game has imperfect information if at least one player does not observe all previous moves when making a decision. Represented by information sets in the game tree.
- Incomplete information** A game has incomplete information if players are uncertain about some aspect of the game, such as other players' payoffs or types. Modelled as a Bayesian game.
- Individually rational payoff** The minimum payoff a player can guarantee regardless of opponents' strategies (the minmax value). In repeated games, the folk theorem applies to payoffs above this threshold.
- Information set** In an extensive-form game, a collection of decision nodes among which a player cannot distinguish. The player must choose the same action at all nodes in the same information set.
- Matching Pennies** A 2x2 zero-sum game where one player wins if outcomes match and the other wins if they differ. It has a unique Nash equilibrium in mixed strategies.
- Mechanism design** The “inverse” of game theory: designing the rules of a game (mechanism) so that self-interested players produce a desired outcome in equilibrium. See Chapter 67.
- Minmax value** The lowest payoff that opponents can force on a player, or equivalently, the highest payoff a player can guarantee against the worst-case opponent behaviour.
- Mixed strategy** A probability distribution over a player's pure strategies. A player randomises among actions according to the specified probabilities. See Chapter 11.
- Monte Carlo method** A computational technique that uses random sampling to estimate quantities or simulate systems. Used throughout Part III for game simulations. See Chapter 33.
- Nash equilibrium (NE)** A strategy profile in which no player can increase their payoff by unilaterally changing their strategy. Every finite game has at least one Nash equilibrium (possibly in mixed strategies). See Chapter 9.
- Normal form (strategic form)** A representation of a simultaneous game as a matrix of payoffs. Each row corresponds to a strategy of one player, each column to a strategy of the other. See Chapter 7.
- Pareto optimal (Pareto efficient)** An outcome is Pareto optimal if no alternative outcome makes at least one player strictly better off without making any player worse off. The Prisoner's Dilemma equilibrium is famously Pareto inefficient.
- Payoff matrix** The matrix encoding all players' payoffs for every combination of strategies in a normal-form game.
- Perfect information** A game has perfect information if every player, when making a decision, observes all previously taken actions. Chess is a canonical example.
- Prisoner's Dilemma** A 2x2 game where each player has a dominant strategy to defect, yet mutual cooperation yields a higher payoff for both. The unique NE is Pareto inefficient.
- Pure strategy** A deterministic plan of action that specifies exactly which action a player will take at every decision point.

- Replicator dynamics** A system of differential equations describing how the proportions of strategies in a population change over time based on their relative fitness. See Chapter 41.
- Shapley value** A unique allocation in cooperative game theory that distributes the total value among players according to their average marginal contribution across all possible orderings. See Chapter 19.
- Signaling game** A dynamic Bayesian game where an informed player (sender) takes an observable action to convey information to an uninformed player (receiver). See Chapter 15.
- Stag Hunt** A 2x2 coordination game with two pure-strategy NE: one payoff-dominant (both hunt Stag) and one risk-dominant (both hunt Hare). It models the tension between cooperation and safety.
- Strategy profile** A specification of one strategy for every player in the game. An n -player game has a strategy profile $(s_1, s_2, \dots, s_n)$.
- Subgame** A portion of an extensive-form game that starts at a singleton information set, contains all successors, and does not break any information set.
- Subgame perfect equilibrium (SPE)** A strategy profile that constitutes a Nash equilibrium in every subgame of the original game. Found by backward induction in games of perfect information. See Chapter 13.
- Support (of a mixed strategy)** The set of pure strategies assigned positive probability in a mixed strategy. At a Nash equilibrium, every strategy in the support must yield the same expected payoff.
- Tit-for-Tat (TFT)** A strategy in repeated games that cooperates in the first round, then copies the opponent's previous action. Famously successful in Axelrod's tournaments. See Chapter 37.
- Zero-sum game** A game in which the sum of all players' payoffs is zero (or constant) for every outcome. One player's gain is exactly another's loss. Matching Pennies is a canonical example.

Bibliography

- Arora, S., Hazan, E., & Kale, S. (2012). The multiplicative weights update method: A meta-algorithm and applications. *Theory of Computing*, 8(1), 121–164. <https://doi.org/10.4086/toc.2012.v008a006>
- Axelrod, R. (1981). The evolution of cooperation. *Science*, 211(4489), 1390–1396. <https://doi.org/10.1126/science.7466396>
- Barasz, M., Christiano, P., Fallenstein, B., Herreshoff, M., LaVictoire, P., & Yudkowsky, E. (2014). Robust cooperation in the prisoner’s dilemma: Program equilibrium via provability logic. *arXiv preprint*.
- Bowling, M., Burch, N., Johanson, M., & Tammelin, O. (2015). Heads-up limit hold’em poker is solved. *Science*, 347(6218), 145–149. <https://doi.org/10.1126/science.1259433>
- Brown, N., & Sandholm, T. (2019). Superhuman ai for multiplayer poker. *Science*, 365(6456), 885–890. <https://doi.org/10.1126/science.aay2400>
- Camerer, C. F., Ho, T.-H., & Chong, J.-K. (2004). A cognitive hierarchy model of games. *Quarterly Journal of Economics*, 119(3), 861–898. <https://doi.org/10.1162/0033553041502225>
- Cesa-Bianchi, N., & Lugosi, G. (2006). *Prediction, learning, and games*. Cambridge University Press.
- Chouldechova, A. (2017). Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2), 153–163. <https://doi.org/10.1089/big.2016.0047>
- Dafoe, A., Hughes, E., Bachrach, Y., Collins, T., McKee, K. R., Leibo, J. Z., Larson, K., & Graepel, T. (2020). Open problems in cooperative ai. *arXiv preprint*.
- Daskalakis, C., Goldberg, P. W., & Papadimitriou, C. H. (2009). The complexity of computing a nash equilibrium. *SIAM Journal on Computing*, 39(1), 195–259. <https://doi.org/10.1137/070699652>
- Dwork, C., Hardt, M., Pitassi, T., Reingold, O., & Zemel, R. (2012). Fairness through awareness. *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, 214–226. <https://doi.org/10.1145/2090236.2090255>
- Fudenberg, D., & Levine, D. K. (1998). *The theory of learning in games*. MIT Press.
- Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., & Bengio, Y. (2014). Generative adversarial nets. *Advances in Neural Information Processing Systems*, 27, 2672–2680.
- Hadfield-Menell, D., Russell, S. J., Abbeel, P., & Dragan, A. (2016). Cooperative inverse reinforcement learning. *Advances in Neural Information Processing Systems*, 29, 3909–3917.
- Hardt, M., Price, E., & Srebro, N. (2016). Equality of opportunity in supervised learning. *Advances in Neural Information Processing Systems*, 29, 3315–3323.
- Hart, S., & Mas-Colell, A. (2000). A simple adaptive procedure leading to correlated equilibrium. *Econometrica*, 68(5), 1127–1150. <https://doi.org/10.1111/1468-0262.00153>
- Hofbauer, J., & Sigmund, K. (1998). Evolutionary games and population dynamics.
- Kennan, J., & Wilson, R. (1993). Bargaining with private information. *Journal of Economic Literature*, 31(1), 45–104.

- Kilbertus, N., Carulla, M. R., Parascandolo, G., Hardt, M., Janzing, D., & Schölkopf, B. (2017). Avoiding discrimination through causal reasoning. *Advances in Neural Information Processing Systems*, 30, 656–666.
- Kleinberg, J., Mullainathan, S., & Raghavan, M. (2017). Inherent trade-offs in the fair determination of risk scores. *Proceedings of Innovations in Theoretical Computer Science*, 43:1–43:23.
- Kuhn, H. W. (1950). A simplified two-person poker. *Contributions to the Theory of Games*, 1, 97–103.
- Lanctot, M., Zambaldi, V., Gruslys, A., Lazaridou, A., Tuyls, K., Pérolat, J., Silver, D., & Graepel, T. (2017). A unified game-theoretic approach to multiagent reinforcement learning. *Advances in Neural Information Processing Systems*, 30, 4190–4203.
- Littman, M. L. (1994). Markov games as a framework for multi-agent reinforcement learning. *Proceedings of the Eleventh International Conference on Machine Learning*, 157–163. <https://doi.org/10.1016/B978-1-55860-335-6.50027-1>
- Lowe, R., Wu, Y., Tamar, A., Harb, J., Abbeel, P., & Mordatch, I. (2017). Multi-agent actor-critic for mixed cooperative-competitive environments. *Advances in Neural Information Processing Systems*, 30, 6379–6390.
- Maynard Smith, J. (1982). *Evolution and the theory of games*. Cambridge University Press.
- Meta Fundamental AI Research Diplomacy Team (FAIR), Bakhtin, A., Brown, N., Dinan, E., Farina, G., Flaherty, C., Fried, D., Goff, A., Gray, J., Hu, H., et al. (2022). Human-level play in the game of diplomacy by combining language models with strategic reasoning. *Science*, 378(6624), 1067–1074. <https://doi.org/10.1126/science.ade9097>
- Moulin, H. (2004). *Fair division and collective welfare*. MIT Press.
- Myerson, R. B. (1991). *Game theory: Analysis of conflict*. Harvard University Press.
- Nagel, R. (1995). Unraveling in guessing games: An experimental study. *American Economic Review*, 85(5), 1313–1326.
- Nash, J. F. (1950). *Equilibrium points in n-person games* (Vol. 36). <https://doi.org/10.1073/pnas.36.1.48>
- Nisan, N., Roughgarden, T., Tardos, É., & Vazirani, V. V. (2007). *Algorithmic game theory*. Cambridge University Press.
- Nowak, M. A., & May, R. M. (1992). Evolutionary games and spatial chaos. *Nature*, 359(6398), 826–829. <https://doi.org/10.1038/359826a0>
- Nowak, M. A., Sasaki, A., Taylor, C., & Fudenberg, D. (2004). Emergence of cooperation and evolutionary stability in finite populations. *Nature*, 428(6983), 646–650. <https://doi.org/10.1038/nature02414>
- Osborne, M. J. (2004). *An introduction to game theory*. Oxford University Press.
- Rawls, J. (1971). *A theory of justice*. Harvard University Press.
- Roth, A. E. (2002). The economist as engineer: Game theory, experimentation, and computation as tools for design economics. *Econometrica*, 70(4), 1341–1378. <https://doi.org/10.1111/1468-0262.00335>
- Roth, A. E., & Sotomayor, M. A. O. (1992). *Two-sided matching: A study in game-theoretic modeling and analysis*. Cambridge University Press.
- Roughgarden, T., & Tardos, É. (2002). How bad is selfish routing? *Journal of the ACM*, 49(2), 236–259. <https://doi.org/10.1145/506147.506153>
- Rubinstein, A. (1982). Perfect equilibrium in a bargaining model. *Econometrica*, 50(1), 97–109. <https://doi.org/10.2307/1912531>
- Santos, F. C., & Pacheco, J. M. (2005). Scale-free networks provide a unifying framework for the emergence of cooperation. *Physical Review Letters*, 95(9), 098104. <https://doi.org/10.1103/PhysRevLett.95.098104>
- Schelling, T. C. (1960). *The strategy of conflict*. Harvard University Press.

- Schuster, P., & Sigmund, K. (1983). Replicator dynamics. *Journal of Theoretical Biology*, 100(3), 533–538. [https://doi.org/10.1016/0022-5193\(83\)90445-9](https://doi.org/10.1016/0022-5193(83)90445-9)
- Sen, A. (1999). *Development as freedom*. Oxford University Press.
- Shapley, L. S. (1953). A value for n-person games. *Annals of Mathematics Studies*, 28, 307–317. <https://doi.org/10.1515/9781400881970-018>
- Shoham, Y., & Leyton-Brown, K. (2009). *Multiagent systems: Algorithmic, game-theoretic, and logical foundations*. Cambridge University Press.
- Silver, D., Huang, A., Maddison, C. J., Guez, A., Sifre, L., van den Driessche, G., Schrittwieser, J., Antonoglou, I., Panneershelvam, V., Lanctot, M., Dieleman, S., Grewe, D., Nham, J., Kalchbrenner, N., Sutskever, I., Lillicrap, T., Leach, M., Kavukcuoglu, K., Graepel, T., & Hassabis, D. (2016). Mastering the game of go with deep neural networks and tree search. *Nature*, 529(7587), 484–489. <https://doi.org/10.1038/nature16961>
- Silver, D., Schrittwieser, J., Simonyan, K., Antonoglou, I., Huang, A., Guez, A., Hubert, T., Baker, L., Lai, M., Bolton, A., Chen, Y., Lillicrap, T., Hui, F., Sifre, L., van den Driessche, G., Graepel, T., & Hassabis, D. (2017). Mastering the game of go without human knowledge. *Nature*, 550(7676), 354–359. <https://doi.org/10.1038/nature24270>
- Soares, N., & Fallenstein, B. (2015). Agent foundations for aligning machine intelligence with human interests: A technical research agenda. *Machine Intelligence Research Institute Technical Report*.
- Stahl, D. O., & Wilson, P. W. (1995). On players’ models of other players: Theory and experimental evidence. *Games and Economic Behavior*, 10(1), 218–254. <https://doi.org/10.1006/game.1995.1031>
- Strang, G. (2022). *Introduction to linear algebra* (6th ed.). Wellesley-Cambridge Press.
- Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction* (2nd ed.). MIT Press.
- Szabó, G., & Fáth, G. (2007). Evolutionary games on graphs. *Physics Reports*, 446(4–6), 97–216. <https://doi.org/10.1016/j.physrep.2007.04.004>
- Taylor, P. D., & Jonker, L. B. (1978). Evolutionary stable strategies and game dynamics. *Mathematical Biosciences*, 40(1–2), 145–156. [https://doi.org/10.1016/0025-5564\(78\)90077-9](https://doi.org/10.1016/0025-5564(78)90077-9)
- Tennenholtz, M. (2004). Program equilibrium. *Games and Economic Behavior*, 49(2), 363–373. <https://doi.org/10.1016/j.geb.2004.02.002>
- Vinyals, O., Babuschkin, I., Czarnecki, W. M., Mathieu, M., Dudzik, A., Chung, J., Choi, D. H., Powell, R., Ewalds, T., Georgiev, P., Oh, J., Horgan, D., Kroiss, M., Danihelka, I., Huang, A., Sifre, L., Cai, T., Agapiou, J. P., Jaderberg, M., ... Silver, D. (2019). Grandmaster level in starcraft ii using multi-agent reinforcement learning. *Nature*, 575(7782), 350–354. <https://doi.org/10.1038/s41586-019-1724-z>
- von Neumann, J., & Morgenstern, O. (1944). *Theory of games and economic behavior*. Princeton University Press.
- Watkins, C. J. C. H., & Dayan, P. (1992). Q-learning. *Machine Learning*, 8(3–4), 279–292. <https://doi.org/10.1007/BF00992698>
- Watts, D. J., & Strogatz, S. H. (1998). Collective dynamics of ‘small-world’ networks. *Nature*, 393(6684), 440–442. <https://doi.org/10.1038/30918>
- Wickham, H., Çetinkaya-Rundel, M., & Grolemund, G. (2023). *R for data science* (2nd ed.). O’Reilly Media.
- Zinkevich, M., Johanson, M., Bowling, M., & Piccione, C. (2007). Regret minimization in games with incomplete information. *Advances in Neural Information Processing Systems*, 20, 1729–1736.