

53. Q-Learning in Games

Learning objectives

- Explain how single-agent Q-learning extends to multi-agent matrix games.
- Implement independent Q-learners for repeated 2-player games in R.
- Demonstrate convergence to Nash equilibrium in coordination games and cycling in zero-sum games.
- Discuss the “rational learning” debate and the limits of independent learning.

53.1. Motivation

Reinforcement learning has achieved remarkable success in single-agent settings: an agent interacts with an environment, receives rewards, and learns a policy that maximizes long-term return. But what happens when the “environment” is another learning agent?

In a matrix game played repeatedly, two Q-learners each try to maximize their own payoff. The environment each agent faces is **non-stationary** — it changes as the other agent updates its policy. This creates a fundamental tension: the convergence guarantees of Q-learning (Chapter 51) rely on a stationary environment, and that assumption breaks in multi-agent settings.

Whether independent Q-learners converge to Nash equilibrium, cycle, or diverge depends on the game’s structure. This chapter explores all three outcomes through simulation.

53.2. Theory

53.2.1. Q-learning recap

In single-agent Q-learning (Sutton & Barto, 2018, Chapter 6), an agent maintains a value estimate $Q(s, a)$ for each state-action pair and updates it after observing a reward r and next state s' :

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] \quad (53.1)$$

For repeated matrix games, there is only one “state” (the game itself), so Q-values reduce to action values $Q(a_i)$ for each of player i ’s actions. The discount factor γ is irrelevant in a stateless setting, and the update simplifies to:

$$Q_i(a_i) \leftarrow Q_i(a_i) + \alpha [r_i - Q_i(a_i)] \quad (53.2)$$

where r_i is the payoff player i received from playing action a_i against the opponent’s (unknown) action.

53.2.2. Independent learners

The simplest multi-agent approach is **independent Q-learning**: each agent runs its own Q-learning algorithm, treating the other agent as part of the environment. No agent models or observes the other's Q-values.

The action selection typically uses an ε -greedy policy: with probability ε the agent explores (chooses a random action), and with probability $1 - \varepsilon$ it exploits (chooses the action with the highest Q-value).

53.2.3. Convergence properties

The behavior of independent Q-learners depends critically on the game:

- **Coordination games** (e.g., Battle of the Sexes): learners tend to converge to one of the pure-strategy Nash equilibria. Which one depends on initial conditions and exploration noise.
- **Zero-sum games** (e.g., Matching Pennies): learners typically **cycle** rather than converge. Each agent's best response shifts the opponent's optimal play, creating a feedback loop.
- **The Prisoner's Dilemma**: learners converge to mutual defection — the unique Nash equilibrium — even though mutual cooperation yields higher payoffs.

Shoham and Leyton-Brown (2009) discuss the “rational learning” debate: under what conditions can independent learners reach equilibrium, and when should we expect them to fail?

53.3. Implementation in R

53.3.1. Q-learning agent

```
q_agent_create <- function(n_actions, alpha = 0.1, epsilon = 0.1) {
  Q <- rep(0, n_actions)
  list(
    choose = function() {
      if (runif(1) < epsilon) {
        sample(n_actions, 1)
      } else {
        ties <- which(Q == max(Q))
        sample(ties, 1)
      }
    },
    update = function(action, reward) {
      Q[action] <- Q[action] + alpha * (reward - Q[action])
    },
    get_Q = function() Q
  )
}
```

53.3.2. Running a repeated game

```
run_q_learning_game <- function(A, B, episodes = 5000,
                                alpha = 0.1, epsilon = 0.1) {
  n1 <- nrow(A)
  n2 <- ncol(A)
  agent1 <- q_agent_create(n1, alpha, epsilon)
  agent2 <- q_agent_create(n2, alpha, epsilon)

  history <- tibble(
    episode = integer(),
    a1 = integer(), a2 = integer(),
    r1 = double(), r2 = double(),
    Q1_1 = double(), Q1_2 = double(),
    Q2_1 = double(), Q2_2 = double()
  )

  for (ep in seq_len(episodes)) {
    a1 <- agent1$choose()
    a2 <- agent2$choose()
    r1 <- A[a1, a2]
    r2 <- B[a1, a2]
    agent1$update(a1, r1)
    agent2$update(a2, r2)

    Q1 <- agent1$get_Q()
    Q2 <- agent2$get_Q()
    history <- bind_rows(history, tibble(
      episode = ep, a1 = a1, a2 = a2, r1 = r1, r2 = r2,
      Q1_1 = Q1[1], Q1_2 = Q1[2],
      Q2_1 = Q2[1], Q2_2 = Q2[2]
    ))
  }
  history
}
```

53.3.3. Coordination game: convergence

```
A_coord <- matrix(c(2, 0, 0, 1), nrow = 2, byrow = TRUE,
                  dimnames = list(c("A", "B"), c("A", "B")))
B_coord <- A_coord

set.seed(42)
hist_coord <- run_q_learning_game(A_coord, B_coord,
                                  episodes = 3000, alpha = 0.1, epsilon = 0.05)
```

```
q_long <- hist_coord |>
  select(episode, Q1_1, Q1_2, Q2_1, Q2_2) |>
```

```

pivot_longer(-episode, names_to = "variable", values_to = "Q") |>
mutate(
  player = if_else(str_starts(variable, "Q1"), "Player 1", "Player 2"),
  action = if_else(str_ends(variable, "_1"), "Action A", "Action B")
)

p_conv <- ggplot(q_long, aes(x = episode, y = Q, colour = action, linetype = player)) +
  geom_line(alpha = 0.7) +
  scale_colour_manual(values = c("Action A" = okabe_ito[1],
                                "Action B" = okabe_ito[2])) +
  theme_publication() +
  labs(title = "Q-Value Convergence: Coordination Game",
       x = "Episode", y = "Q-value",
       colour = "Action", linetype = "Player")

p_conv
save_pub_fig(p_conv, "q-convergence-coordination", width = 7, height = 5)

```




Figure 53.1.: Q-value evolution in a coordination game. Both agents converge to action A (the payoff-dominant Nash equilibrium), with Q-values stabilizing near the equilibrium payoff of 2.

53.3.4. Zero-sum game: cycling

```

# Matching Pennies
A_mp <- matrix(c(1, -1, -1, 1), nrow = 2, byrow = TRUE,
              dimnames = list(c("H", "T"), c("H", "T")))
B_mp <- -A_mp

```

```

set.seed(42)
hist_mp <- run_q_learning_game(A_mp, B_mp,
                              episodes = 3000, alpha = 0.05, epsilon = 0.1)

q_long_mp <- hist_mp |>
  select(episode, Q1_1, Q1_2, Q2_1, Q2_2) |>
  pivot_longer(-episode, names_to = "variable", values_to = "Q") |>
  mutate(
    player = if_else(str_starts(variable, "Q1"), "Player 1", "Player 2"),
    action = if_else(str_ends(variable, "_1"), "Heads", "Tails")
  )

p_cycle <- ggplot(q_long_mp, aes(x = episode, y = Q, colour = action, linetype = player)) +
  geom_line(alpha = 0.7) +
  scale_colour_manual(values = c("Heads" = okabe_ito[1],
                                "Tails" = okabe_ito[2])) +
  theme_publication() +
  labs(title = "Q-Value Cycling: Matching Pennies",
       x = "Episode", y = "Q-value",
       colour = "Action", linetype = "Player")

p_cycle
save_pub_fig(p_cycle, "q-cycling-matching-pennies", width = 7, height = 5)

```




Figure 53.2.: Q-value evolution in Matching Pennies (zero-sum). Instead of converging, Q-values oscillate as each agent's adaptation shifts the other's optimal play. The mixed NE at (0.5, 0.5) is never stably reached.

53.4. Worked example

Let us trace the first 20 episodes of the coordination game in detail:

```
trace <- hist_coord |>
  filter(episode <= 20) |>
  mutate(
    a1_label = if_else(a1 == 1, "A", "B"),
    a2_label = if_else(a2 == 1, "A", "B"),
    coordinated = a1 == a2
  )

cat("First 20 episodes of coordination game Q-learning:\n\n")
```

```
#> First 20 episodes of coordination game Q-learning:
```

```
trace |>
  select(episode, a1_label, a2_label, r1, r2, coordinated) |>
  rename(P1 = a1_label, P2 = a2_label, `P1 payoff` = r1,
    `P2 payoff` = r2, Coordinated = coordinated) |>
  print(n = 20)
```

```
#> # A tibble: 20 x 6
#>   episode P1    P2    `P1 payoff` `P2 payoff` Coordinated
#>   <int> <chr> <chr>         <dbl>         <dbl> <lgl>
#> 1     1     1 A     A             2             2 TRUE
#> 2     2     2 A     A             2             2 TRUE
#> 3     3     3 A     A             2             2 TRUE
#> 4     4     4 A     A             2             2 TRUE
#> 5     5     5 A     A             2             2 TRUE
#> 6     6     6 A     A             2             2 TRUE
#> 7     7     7 A     A             2             2 TRUE
#> 8     8     8 A     A             2             2 TRUE
#> 9     9     9 A     B             0             0 FALSE
#> 10    10    10 B     A             0             0 FALSE
#> 11    11    11 A     B             0             0 FALSE
#> 12    12    12 A     A             2             2 TRUE
#> 13    13    13 A     A             2             2 TRUE
#> 14    14    14 A     B             0             0 FALSE
#> 15    15    15 A     A             2             2 TRUE
#> 16    16    16 A     A             2             2 TRUE
#> 17    17    17 A     A             2             2 TRUE
#> 18    18    18 A     A             2             2 TRUE
#> 19    19    19 A     A             2             2 TRUE
#> 20    20    20 A     A             2             2 TRUE
```

```
cat(sprintf("\nCoordination rate (first 20): %.0f%%\n",
  100 * mean(trace$coordinated)))
```

```
#>
#> Coordination rate (first 20): 80%

cat(sprintf("Coordination rate (last 500): %.0f%%\n",
           100 * mean((hist_coord |> filter(episode > 2500))$a1 ==
                     (hist_coord |> filter(episode > 2500))$a2)))

#> Coordination rate (last 500): 94%
```

Early episodes show exploration: agents try both actions and sometimes miscoordinate. As Q-values diverge (one action becomes clearly better), the ε -greedy policy increasingly selects the high-Q action, and coordination improves. By the final 500 episodes, coordination is near-perfect.

In Matching Pennies, by contrast, no such convergence occurs — the Q-values never stabilize because each agent’s adaptation continually shifts the landscape the other agent faces.

53.5. Extensions

- **Decaying exploration:** Using $\varepsilon_t = \varepsilon_0/t$ guarantees convergence in stationary environments but can lock in suboptimal play in non-stationary multi-agent settings.
- **Win-or-Learn-Fast (WoLF):** Varies the learning rate based on whether the agent is winning or losing, achieving convergence in a broader class of games.
- **Multi-agent RL** (Chapter 55) covers advanced approaches: MADDPG, PSRO, and the challenge of non-stationarity at scale.
- **Reinforcement learning foundations** (Chapter 51) provides the single-agent background that this chapter builds on.

For the theoretical debate on rational learning in games, see Shoham and Leyton-Brown (2009) (Chapter 7).

Exercises

1. **Prisoner’s Dilemma.** Run the Q-learning simulation on the standard Prisoner’s Dilemma ($T = 5$, $R = 3$, $P = 1$, $S = 0$). Do both agents converge to mutual defection? How many episodes does it take? Plot the Q-value evolution.
2. **Learning rate sensitivity.** Run the coordination game with $\alpha \in \{0.01, 0.1, 0.5\}$ and fixed $\varepsilon = 0.05$. How does the learning rate affect convergence speed and stability? Plot Q-value trajectories for all three.
3. **Boltzmann exploration.** Replace ε -greedy action selection with Boltzmann (softmax) exploration: $P(a_i) = \exp(Q(a_i)/\tau) / \sum_{a'} \exp(Q(a')/\tau)$, where τ is a temperature parameter. Implement this and compare performance in the coordination game with ε -greedy.

Solutions appear in Appendix H.

