

Cooperative Game Theory

An introduction to cooperative game theory with transferable utility. Covers the characteristic function, the core, and the Shapley value with computational examples in R, including a worked cost-sharing problem for three municipalities.

Learning objectives

- Define a transferable-utility (TU) game using its characteristic function and explain how it differs from non-cooperative models.
- Compute the core of a cooperative game and determine whether a given allocation lies in the core.
- Calculate the Shapley value for small games by hand and in R, and interpret its fairness properties.
- Apply cooperative game theory to practical cost-sharing and voting-power problems.

Motivation

Non-cooperative game theory, as developed in @ref(sec-nash-equilibrium) and preceding chapters, models strategic interaction by focusing on individual players' strategies and the equilibria that emerge. But many real-world situations are better described by asking which **coalitions** can form and how the resulting surplus should be divided. When three municipalities consider building a shared water treatment plant, the key questions are not about individual strategic moves but about which subsets of towns benefit from cooperating and how the cost savings should be split.

Cooperative game theory, whose foundations were laid by @von-neumann1944, takes coalitions as the primitive unit of analysis. Instead of specifying each player's strategy set, a cooperative game specifies the **worth** (or value) that each possible coalition can generate. The central solution concepts — the **core** and the **Shapley value** — then determine which divisions of the total payoff are stable against defection and which are uniquely fair.

This chapter focuses on **transferable-utility (TU) games**, where utility can be freely redistributed among coalition members (e.g., via monetary side payments). We introduce the characteristic function, define stability and fairness concepts, and implement everything computationally in R. For R packages that automate many of these calculations, see @ref(sec-coopgame-package) and @ref(sec-gametheory-package).

Theory

Characteristic function games

A **TU cooperative game** is a pair (N, v) where:

- $N = \{1, 2, \dots, n\}$ is the set of players.
- $v : 2^N \rightarrow \mathbb{R}$ is the **characteristic function** assigning a value $v(S)$ to each coalition $S \subseteq N$, with $v(\emptyset) = 0$.

The value $v(S)$ represents the total payoff that members of S can guarantee themselves by cooperating, regardless of what players outside S do. A game is **superadditive** if $v(S \cup T) \geq v(S) + v(T)$ for all disjoint $S, T \subseteq N$, meaning that merging coalitions never destroys value.

Imputations

An **imputation** is a payoff vector $x = (x_1, \dots, x_n)$ satisfying:

1. **Efficiency:** $\sum_{i \in N} x_i = v(N)$.
2. **Individual rationality:** $x_i \geq v(\{i\})$ for all $i \in N$.

Efficiency requires that the grand coalition's total value is fully distributed. Individual rationality ensures no player receives less than they could obtain alone.

The core

Definition: The Core

The **core** of a TU game (N, v) is the set of imputations x such that no coalition has an incentive to deviate:

$$\sum_{i \in S} x_i \geq v(S) \quad \text{for all } S \subseteq N (\#eq : \text{core} - def) \quad (1)$$

An allocation in the core is **stable** in the sense that no subset of players can do better by breaking away from the grand coalition. The core may be empty (some games have no stable allocation), a single point, or a convex polytope.

The Shapley value

While the core captures stability, it does not provide a unique solution. The **Shapley value** provides a unique allocation based on fairness axioms.

Definition: Shapley Value

The Shapley value of player i in a TU game (N, v) is:

$$\phi_i(v) = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!(n - |S| - 1)!}{n!} [v(S \cup \{i\}) - v(S)] (\#eq : \text{shapley} - value) \quad (2)$$

The term $v(S \cup \{i\}) - v(S)$ is player i 's **marginal contribution** to coalition S . The Shapley value averages this marginal contribution over all possible orderings in which the grand coalition could be assembled. It is the unique allocation satisfying four axioms: efficiency, symmetry, additivity, and the null-player property.

As @osborne2004 [Chapter 18] emphasizes, the Shapley value can be interpreted as the expected marginal contribution of a player when players join the coalition in a uniformly random order.

Implementation in R

Defining a cooperative game

We consider a three-player cost-sharing game. Three municipalities (A, B, C) can build a water treatment plant individually or in coalition. The costs are:

Coalition	Cost	Savings $v(S)$
{A}	100	0

Coalition	Cost	Savings $v(S)$
{B}	140	0
{C}	120	0
{A, B}	200	40
{A, C}	180	40
{B, C}	220	40
{A, B, C}	300	60

The characteristic function $v(S)$ measures cost *savings* relative to individual construction. We normalize $v(\{i\}) = 0$ for each individual player.

```
# Define the 3-player TU game (cost savings)
players <- c("A", "B", "C")
n <- length(players)

# Characteristic function: value of each coalition
# Using binary encoding: A=1, B=2, C=4
v <- c(
  "0" = 0,      # empty

  "A" = 0,      # {A}
  "B" = 0,      # {B}
  "C" = 0,      # {C}
  "AB" = 40,    # {A, B}
  "AC" = 40,    # {A, C}
  "BC" = 40,    # {B, C}
  "ABC" = 60    # {A, B, C}
)

cat("Characteristic function v(S):\n")
```

```
#> Characteristic function v(S):
```

```
for (name in names(v)) {
  if (name != "0") {
    cat(sprintf("  v({%s}) = %d\n",
               paste(strsplit(name, "")[[1]], collapse = ", "), v[name]))
  }
}
```

```
#>  v({A}) = 0
#>  v({B}) = 0
#>  v({C}) = 0
#>  v({A, B}) = 40
#>  v({A, C}) = 40
#>  v({B, C}) = 40
#>  v({A, B, C}) = 60
```

Computing the Shapley value

```

# Permutation generator (no external dependencies)
permn <- function(x) {
  if (length(x) <= 1) return(list(x))
  result <- list()
  for (i in seq_along(x)) {
    rest <- permn(x[-i])
    for (p in rest) result <- c(result, list(c(x[i], p)))
  }
  result
}

# Compute Shapley value by enumeration of permutations
compute_shapley <- function(players, v_func) {
  n <- length(players)
  perms <- permn(players)
  marginals <- matrix(0, nrow = length(perms), ncol = n,
    dimnames = list(NULL, players))

  for (k in seq_along(perms)) {
    perm <- perms[[k]]
    for (i in seq_along(perm)) {
      player <- perm[i]
      if (i == 1) {
        predecessors <- character(0)
      } else {
        predecessors <- perm[1:(i - 1)]
      }
      coal_with <- paste(sort(c(predecessors, player)), collapse = "")
      coal_without <- paste(sort(predecessors), collapse = "")
      if (coal_without == "") coal_without <- "0"
      marginals[k, player] <- v_func[coal_with] - v_func[coal_without]
    }
  }

  shapley_vals <- colMeans(marginals)
  list(values = shapley_vals, marginals = marginals)
}

shapley_result <- compute_shapley(players, v)
shapley_vals <- shapley_result$values

cat("Shapley values:\n")

```

```
#> Shapley values:
```

```

for (p in players) {
  cat(sprintf("  phi(%s) = %.4f\n", p, shapley_vals[p]))
}

```

```

#>  phi(A) = 20.0000
#>  phi(B) = 20.0000
#>  phi(C) = 20.0000

```

```
cat(sprintf("\nSum of Shapley values: %.4f (should equal v(N) = %d)\n",
           sum(shapley_vals), v["ABC"]))
```

```
#>
```

```
#> Sum of Shapley values: 60.0000 (should equal v(N) = 60)
```

Checking the core

```
# An allocation is in the core if it satisfies all coalition constraints
```

```
check_core <- function(x, players, v_func) {
  n <- length(players)
  violations <- character(0)

  # Check all non-empty subsets
  for (k in 1:(2^n - 1)) {
    bits <- as.integer(intToBits(k)[1:n])
    coalition <- players[bits == 1]
    coal_key <- paste(sort(coalition), collapse = "")
    coal_val <- v_func[coal_key]
    alloc_sum <- sum(x[coalition])

    if (alloc_sum < coal_val - 1e-10) {
      violations <- c(violations,
                     sprintf("{%s}: sum = %.2f < v = %.0f",
                             paste(coalition, collapse = ", "),
                             alloc_sum, coal_val))
    }
  }

  list(in_core = length(violations) == 0, violations = violations)
}
```

```
# Check if Shapley value is in the core
```

```
core_result <- check_core(shapley_vals, players, v)
cat("Is the Shapley value in the core? ",
    ifelse(core_result$in_core, "YES", "NO"), "\n")
```

```
#> Is the Shapley value in the core? YES
```

```
if (!core_result$in_core) {
  cat("Violations:\n")
  for (viol in core_result$violations) cat(" ", viol, "\n")
}
```

```
# Display all coalition constraints
```

```
cat("\nCore constraints:\n")
```

```
#>
```

```
#> Core constraints:
```

```

coalitions <- list(
  c("A"), c("B"), c("C"),
  c("A", "B"), c("A", "C"), c("B", "C")
)

for (coal in coalitions) {
  coal_key <- paste(sort(coal), collapse = "")
  cat(sprintf("  x(%s) = %.2f >= v({%s}) = %d : %s\n",
    paste(coal, collapse = " + "),
    sum(shapley_vals[coal]),
    paste(coal, collapse = ", "),
    v[coal_key],
    ifelse(sum(shapley_vals[coal]) >= v[coal_key] - 1e-10,
      "satisfied", "VIOLATED")))
}

#>  x(A) = 20.00 >= v({A}) = 0 : satisfied
#>  x(B) = 20.00 >= v({B}) = 0 : satisfied
#>  x(C) = 20.00 >= v({C}) = 0 : satisfied
#>  x(A + B) = 40.00 >= v({A, B}) = 40 : satisfied
#>  x(A + C) = 40.00 >= v({A, C}) = 40 : satisfied
#>  x(B + C) = 40.00 >= v({B, C}) = 40 : satisfied

```

Publication figure: Shapley value comparison

```

shapley_df <- tibble(
  player = factor(players, levels = players),
  shapley = shapley_vals[players]
)

p_shapley <- ggplot(shapley_df, aes(x = player, y = shapley, fill = player)) +
  geom_col(width = 0.6, show.legend = FALSE) +
  geom_text(aes(label = sprintf("%.1f", shapley)),
    vjust = -0.5, size = 4, fontface = "bold") +
  scale_fill_manual(values = okabe_ito[1:3]) +
  scale_y_continuous(name = "Shapley value (cost savings)",
    limits = c(0, max(shapley_vals) * 1.3),
    expand = expansion(mult = c(0, 0.05))) +
  scale_x_discrete(name = "Municipality") +
  labs(title = "Shapley Value Allocation: Water Treatment Cost Sharing") +
  theme_publication()

p_shapley

save_pub_fig(p_shapley, "shapley-values-3player")

```

Marginal contributions table




Figure 1: Shapley values for the three-player water treatment cost-sharing game. Each bar represents a municipality's fair share of the total cost savings (60 units) based on average marginal contributions across all coalition formation orders. The Shapley value assigns equal shares (20 each) due to the symmetric structure of pair-wise savings in this game.

```
# Show marginal contributions in each ordering
perms <- permn(players)

mc_table <- tibble(
  ordering = sapply(perms, paste, collapse = " -> "),
  mc_A = shapley_result$marginals[, "A"],
  mc_B = shapley_result$marginals[, "B"],
  mc_C = shapley_result$marginals[, "C"]
)

mc_table |>
  gt() |>
  tab_header(
    title = "Marginal Contributions by Player Ordering",
    subtitle = "Each row shows one permutation and the marginal contribution of each player when joining"
  ) |>
  cols_label(
    ordering = "Ordering",
    mc_A = "MC(A)",
    mc_B = "MC(B)",
    mc_C = "MC(C)"
  ) |>
  cols_align(align = "center", columns = starts_with("mc")) |>
  grand_summary_rows(
    fns = list(Average = ~ mean(.)),
    columns = starts_with("mc"),
```

Marginal Contributions by Player Ordering

Each row shows one permutation and the marginal contribution of each player when joining in that order

	Ordering	MC(A)	MC(B)	MC(C)
	A -> B -> C	0	40	20
	A -> C -> B	0	20	40
	B -> A -> C	40	0	20
	B -> C -> A	20	0	40
	C -> A -> B	40	20	0
	C -> B -> A	20	40	0
Average	—	20.00	20.00	20.00

```
fmt = ~ fmt_number(., decimals = 2)
)
```

Worked example

We walk through the Shapley value calculation for municipality A in the water treatment cost-sharing game.

Step 1: List all permutations. With three players, there are $3! = 6$ orderings.

Step 2: Compute A's marginal contribution in each ordering.

- **A-B-C:** A joins the empty coalition. $v(\{A\}) - v(\emptyset) = 0 - 0 = 0$.
- **A-C-B:** Same as above. $MC(A) = 0$.
- **B-A-C:** A joins $\{B\}$. $v(\{A, B\}) - v(\{B\}) = 40 - 0 = 40$.
- **C-A-B:** A joins $\{C\}$. $v(\{A, C\}) - v(\{C\}) = 40 - 0 = 40$.
- **B-C-A:** A joins $\{B, C\}$. $v(\{A, B, C\}) - v(\{B, C\}) = 60 - 40 = 20$.
- **C-B-A:** A joins $\{B, C\}$. Same as above. $MC(A) = 20$.

Step 3: Average the marginal contributions.

$$\phi_A = \frac{0 + 0 + 40 + 40 + 20 + 20}{6} = \frac{120}{6} = 20$$

Step 4: Verify efficiency. By symmetry of the game (all two-player coalitions have the same value), $\phi_A = \phi_B = \phi_C = 20$. The sum is $20 + 20 + 20 = 60 = v(N)$, confirming efficiency.

```
# Verify step by step for player A
mc_A <- c(
  v["A"] - v["0"],      # A first
  v["A"] - v["0"],      # A first (other order)
  v["AB"] - v["B"],     # A joins B
  v["AC"] - v["C"],     # A joins C
  v["ABC"] - v["BC"],   # A joins BC
  v["ABC"] - v["BC"]    # A joins BC (other order)
)
```


Final Cost Allocation via Shapley Value

Municipality	Stand-alone cost	Shapley share of savings	Cost after sharing
A	100	20.0	80.0
B	140	20.0	120.0
C	120	20.0	100.0

```
cat("Player A marginal contributions:", mc_A, "\n")
```

```
#> Player A marginal contributions: 0 0 40 40 20 20
```

```
cat(sprintf("Shapley value for A: %.0f / 6 = %.4f\n",
            sum(mc_A), mean(mc_A)))
```

```
#> Shapley value for A: 120 / 6 = 20.0000
```

```
# Summary
final_allocation <- tibble(
  Municipality = players,
  `Stand-alone cost` = c(100, 140, 120),
  `Shapley share of savings` = shapley_vals[players],
  `Cost after sharing` = c(100, 140, 120) - shapley_vals[players]
)

final_allocation |>
  gt() |>
  tab_header(
    title = "Final Cost Allocation via Shapley Value"
  ) |>
  fmt_number(columns = c(`Shapley share of savings`, `Cost after sharing`),
             decimals = 1) |>
  cols_align(align = "center")
```

Step 5: Interpret. Each municipality saves 20 units from cooperation. Municipality A's cost drops from 100 to 80, B's from 140 to 120, and C's from 120 to 100. The Shapley value splits the 60 units of savings equally because the game's structure treats all pair-wise coalitions symmetrically. In an asymmetric game — say, where {A,B} saves more than {A,C} — the Shapley value would allocate more savings to the players whose participation creates more value.

Extensions

Cooperative game theory extends in several important directions:

- **Non-transferable utility (NTU) games** drop the assumption that utility can be freely redistributed. The solution concepts become more complex; the NTU Shapley value and the Nash bargaining solution are commonly used.

- **Voting games** are a special class of TU games where each coalition either wins (value 1) or loses (value 0). The Shapley–Shubik power index measures each voter’s influence and is widely applied in political science.
- **Cost allocation** problems in operations research use the Shapley value and the nucleolus (which minimizes the maximum dissatisfaction of any coalition) to allocate shared infrastructure costs.
- **Computational complexity.** Computing the Shapley value requires summing over all 2^n coalitions (or all $n!$ permutations), which becomes infeasible for large n . Approximation algorithms based on random sampling of permutations are used in practice; see @ref(sec-coopgame-package) and @ref(sec-gametheory-package) for R implementations.

For the mathematical foundations of cooperative game theory, see @von-neumann1944 and @osborne2004 [Chapters 16–18]. A computational perspective is given in @shoham2009.

Exercises

1. **Asymmetric savings.** Modify the water treatment game so that $v(\{A, B\}) = 50$, $v(\{A, C\}) = 30$, $v(\{B, C\}) = 40$, and $v(\{A, B, C\}) = 80$. Compute the Shapley value for each player. Which municipality benefits most from cooperation?
2. **Core membership.** For the modified game in Exercise 1, determine whether the Shapley value lies in the core. If not, find an allocation that does lie in the core, or prove that the core is empty.
3. **Voting power.** Consider a weighted voting game with four players and weights $(4, 3, 2, 1)$ and a quota of 6 (a coalition wins if its total weight is at least 6). Compute the Shapley–Shubik power index for each player. Compare it to each player’s weight share. (*Hint: enumerate all $4! = 24$ orderings and identify each player’s pivot position.*)
4. **Glove game.** In the “glove game,” players 1 and 2 each have a left glove and player 3 has a right glove. A pair of gloves (one left, one right) is worth 1; individual gloves are worthless. Write the characteristic function, compute the Shapley value, and find the core. Is the Shapley value in the core?
5. **Scaling to larger games.** Implement a Monte Carlo approximation of the Shapley value that samples K random permutations instead of enumerating all $n!$. Test it on a 10-player game where $v(S) = |S|^2$ and compare the approximation error for $K = 100$, $K = 1000$, and $K = 10000$.

Solutions appear in @ref(sec-solutions).