

# Spatial Prisoner's Dilemma

A computational exploration of Nowak and May's spatial Prisoner's Dilemma, showing how local interaction on a lattice allows cooperation to survive through cluster formation — even without memory, reciprocity, or reputation.

## Learning objectives

- Explain why spatial structure can sustain cooperation in the Prisoner's Dilemma without reciprocity or punishment.
- Implement a lattice-based spatial PD with Moore neighborhood interaction and imitation dynamics in R.
- Visualize the emergence and persistence of cooperator clusters on a grid.
- Analyze how the fraction of cooperators evolves over time under different payoff parameters.

## Motivation

In @ref(sec-axelrod-tournament), cooperation emerged because players could remember past interactions and condition their behavior on history. But many biological and social systems lack this luxury. Bacteria, plants, and even commuters on a highway interact with nearby neighbors without tracking individual histories. Can cooperation survive when there is no memory, no reputation, and no punishment — only spatial structure?

In 1992, Martin Nowak and Robert May [-@nowak1992] showed that the answer is yes. They placed players on a two-dimensional grid where each cell plays a one-shot Prisoner's Dilemma with its immediate neighbors. After each round, every cell adopts the strategy of whichever neighbor (including itself) scored highest. The result was striking: cooperators formed compact clusters that could resist invasion by defectors, producing intricate, dynamic spatial patterns.

This chapter replicates their model in R. We will build a lattice simulation, watch cooperation clusters emerge, and measure how population-level cooperation rates evolve over time.

## Theory

### From well-mixed to spatial populations

Standard evolutionary game theory assumes a **well-mixed population**: every individual is equally likely to interact with every other. Under this assumption, defection dominates in a one-shot Prisoner's Dilemma because a defector always earns more than a cooperator against any given opponent.

Spatial structure breaks this assumption. When interaction is **local** — each player meets only its geographic neighbors — cooperators can form clusters. Within a cluster, cooperators interact mainly with other cooperators, earning the mutual cooperation payoff  $R$  many times. A defector on the boundary of such a cluster exploits its cooperator neighbors but has fewer cooperator neighbors than cooperators in the cluster interior do.

### The Nowak–May model

The model works as follows:

1. **Grid**: An  $N \times N$  lattice with periodic boundary conditions (a torus, so edges wrap around).
2. **Strategies**: Each cell is either a cooperator (C) or a defector (D).

3. **Interaction:** Each cell plays a one-shot PD with each of its eight Moore neighbors and itself (nine games total). The cell's score is the sum of payoffs from all nine games.
4. **Update:** Synchronously, every cell adopts the strategy of the highest-scoring cell in its Moore neighborhood (including itself). Ties are broken in favor of the current occupant.

With canonical PD payoffs ( $T > R > P > S$ ), the dynamics depend critically on the temptation-to-defect parameter  $b = T/R$ . For moderate values of  $b$ , cooperators and defectors coexist in dynamic spatial patterns. For very high  $b$ , defection takes over; for low  $b$ , cooperation dominates.

### Why clusters protect cooperators

Consider a  $3 \times 3$  block of cooperators surrounded by defectors. The cooperator at the center plays nine cooperators (including itself), scoring  $9R$ . A defector adjacent to this block has at most three cooperator neighbors, scoring at most  $3T + 6P$ . When  $9R > 3T + 6P$  — that is, when  $b < 3(R - P/R) + P$  approximately — the interior cooperator outscores the boundary defector, and the cluster holds. This is the geometric logic behind spatial cooperation.

## Implementation in R

### Grid initialization and neighbor scoring

```
initialize_grid <- function(n, prob_coop = 0.5) {
  matrix(sample(c("C", "D"), n * n, replace = TRUE,
    prob = c(prob_coop, 1 - prob_coop)),
    nrow = n, ncol = n)
}

get_moore_neighbors <- function(grid, i, j) {
  n <- nrow(grid)
  rows <- ((i - 2):(i)) %% n + 1
  cols <- ((j - 2):(j)) %% n + 1
  grid[rows, cols]
}

score_cell <- function(grid, i, j) {
  neighbors <- get_moore_neighbors(grid, i, j)
  focal <- grid[i, j]
  total <- 0
  for (s in as.vector(neighbors)) {
    total <- total + play_round(focal, s)[1]
  }
  total
}

compute_scores <- function(grid) {
  n <- nrow(grid)
  scores <- matrix(0, nrow = n, ncol = n)
  for (i in seq_len(n)) {
    for (j in seq_len(n)) {
      scores[i, j] <- score_cell(grid, i, j)
    }
  }
}
```

```

}
scores
}

```

Update rule: imitate the best neighbor

```

update_grid <- function(grid, scores) {
  n <- nrow(grid)
  new_grid <- grid
  for (i in seq_len(n)) {
    for (j in seq_len(n)) {
      rows <- ((i - 2):(i)) %% n + 1
      cols <- ((j - 2):(j)) %% n + 1
      neighborhood_scores <- scores[rows, cols]
      best <- which(neighborhood_scores == max(neighborhood_scores),
                    arr.ind = TRUE)
      if (nrow(best) == 1) {
        bi <- rows[best[1, 1]]
        bj <- cols[best[1, 2]]
        new_grid[i, j] <- grid[bi, bj]
      }
      # Tie: keep current strategy (new_grid[i,j] already equals grid[i,j])
    }
  }
  new_grid
}

```

Running the simulation

```

run_spatial_pd <- function(n, generations, probab_coop = 0.5, seed = 42) {
  set.seed(seed)
  grid <- initialize_grid(n, probab_coop)

  history <- vector("list", generations + 1)
  history[[1]] <- grid
  coop_frac <- numeric(generations + 1)
  coop_frac[1] <- mean(grid == "C")

  for (g in seq_len(generations)) {
    scores <- compute_scores(grid)
    grid <- update_grid(grid, scores)
    history[[g + 1]] <- grid
    coop_frac[g + 1] <- mean(grid == "C")
  }

  list(history = history, coop_frac = coop_frac)
}

```

## Visualizing a grid snapshot

```
grid_to_df <- function(grid, gen) {  
  n <- nrow(grid)  
  expand.grid(row = seq_len(n), col = seq_len(n)) |>  
    mutate(strategy = as.vector(grid),  
           generation = gen)  
}
```

## Worked example

We simulate a  $30 \times 30$  grid with 50% initial cooperators for 100 generations, using the canonical PD payoffs ( $T = 5$ ,  $R = 3$ ,  $P = 1$ ,  $S = 0$ ).

```
result <- run_spatial_pd(n = 30, generations = 100, prob_coop = 0.5, seed = 42)
```

```
snapshot_gens <- c(0, 5, 20, 100)  
snapshot_df <- map_dfr(snapshot_gens, function(g) {  
  grid_to_df(result$history[[g + 1]], g)  
})  
snapshot_df$generation <- factor(  
  snapshot_df$generation,  
  levels = snapshot_gens,  
  labels = paste("Generation", snapshot_gens)  
)  
  
p_snapshots <- ggplot(snapshot_df,  
                      aes(x = col, y = row, fill = strategy)) +  
  geom_tile() +  
  facet_wrap(~generation, ncol = 2) +  
  scale_fill_manual(values = c("C" = okabe_ito[1], "D" = okabe_ito[5]),  
                   labels = c("C" = "Cooperate", "D" = "Defect"),  
                   name = "Strategy") +  
  coord_equal() +  
  theme_publication() +  
  theme(axis.text = element_blank(),  
        axis.ticks = element_blank(),  
        panel.grid.major = element_blank()) +  
  labs(title = "Spatial Prisoner's Dilemma: Grid Snapshots",  
       x = NULL, y = NULL)  
  
p_snapshots
```

```
save_pub_fig(p_snapshots, "spatial-pd-snapshots", width = 8, height = 8)
```

The initial random arrangement (generation 0) shows a salt-and-pepper mixture. By generation 5, cooperators have begun to aggregate into clusters while isolated cooperators surrounded by defectors have been converted. By generation 20, the pattern has largely stabilized: compact cooperator clusters persist, separated by defector boundaries. The clusters are not static — their edges shift over time — but the overall cooperation level remains roughly stable.

### Spatial Prisoner's Dilemma: Grid Snapshots




Figure 1: Grid snapshots at generations 0, 5, 20, and 100. Cooperators (orange) form compact clusters that resist invasion by defectors (blue). By generation 20, the spatial pattern has largely stabilized into a dynamic equilibrium.

```
coop_df <- tibble(
  generation = 0:100,
  cooperation = result$coop_frac
)

p_coop <- ggplot(coop_df, aes(x = generation, y = cooperation)) +
  geom_line(colour = okabe_ito[1], linewidth = 1) +
  geom_hline(yintercept = mean(result$coop_frac[51:101]),
    linetype = "dashed", colour = okabe_ito[8]) +
  scale_y_continuous(labels = label_percent(), limits = c(0, 1)) +
  theme_publication() +
  labs(title = "Cooperation Fraction Over Time",
    x = "Generation", y = "Fraction cooperating")

p_coop
```




Figure 2: Cooperation fraction over 100 generations. After an initial transient where isolated cooperators are eliminated, the cooperation rate stabilizes around a dynamic equilibrium. Spatial structure sustains a substantial level of cooperation even in the one-shot Prisoner's Dilemma.

```
save_pub_fig(p_coop, "spatial-pd-cooperation", width = 7, height = 5)
```

```
cat(sprintf("Initial cooperation rate: %.1f%%\n",
  result$coop_frac[1] * 100))
```

```
#> Initial cooperation rate: 47.8%
```

```
cat(sprintf("Final cooperation rate:  %.1f%%\n",
           result$coop_frac[101] * 100))
```

```
#> Final cooperation rate:  5.1%
```

```
cat(sprintf("Mean cooperation (gen 50-100): %.1f%%\n",
           mean(result$coop_frac[51:101]) * 100))
```

```
#> Mean cooperation (gen 50-100): 5.1%
```

The cooperation rate drops sharply in the first few generations as isolated cooperators are eliminated. It then stabilizes as surviving cooperators form self-reinforcing clusters. The long-run cooperation rate depends on the payoff parameters — with the canonical values used here, spatial structure sustains cooperation at a level far above zero, even though defection would dominate in a well-mixed population.

## Extensions

- **Varying temptation:** As the temptation payoff  $T$  increases (holding other payoffs fixed), the sustainable cooperation rate decreases. There is a critical threshold above which cooperators go extinct even with spatial structure. Experimenting with this threshold is a good exercise.
- **Stochastic update rules:** Instead of deterministic imitation of the best neighbor, one can use probabilistic rules where the probability of adopting a neighbor's strategy increases with the neighbor's payoff advantage. This connects to the **Fermi update rule** and reduces sensitivity to small payoff differences.
- **Network structure:** The lattice is a special case of a network. Cooperation dynamics on random graphs, scale-free networks, and small-world networks exhibit qualitatively different behavior ([@ref\(sec-network-games\)](#)).
- **Evolutionary stability:** The spatial model shows that cooperation can be an evolutionarily stable configuration even when it is not an ESS in the well-mixed case ([@ref\(sec-ess\)](#)).
- For the original model, see [@nowak1992](#). For a comprehensive review of spatial evolutionary games, see [@szabo2007](#).

## Exercises

1. **Effect of grid size.** Run the spatial PD on grids of size  $10 \times 10$ ,  $30 \times 30$ , and  $50 \times 50$  with the same initial cooperation probability (50%). How does the long-run cooperation rate depend on grid size? Explain why boundary effects matter more on smaller grids.
2. **Temptation threshold.** Fix  $R = 3$ ,  $P = 1$ ,  $S = 0$  and vary  $T$  from 3.5 to 8 in increments of 0.5. For each value, run the simulation for 100 generations and record the final cooperation rate. Plot cooperation rate vs.  $T$  and identify the approximate threshold where cooperators go extinct. (*Hint:* modify `play_round()` to accept custom payoffs.)
3. **Initial cluster.** Instead of random initial placement, start with a single  $5 \times 5$  block of cooperators in the center of a  $30 \times 30$  grid of defectors. Does the cooperator block survive? Grow? How does the outcome depend on  $T$ ?

Solutions appear in [@ref\(sec-solutions\)](#).