

Reinforcement Learning Foundations

Markov decision processes, the Bellman equation, tabular Q-learning and SARSA from scratch, and the exploration-exploitation trade-off via epsilon-greedy and UCB.

Learning objectives

- Define a Markov decision process (MDP) and derive the Bellman optimality equation.
- Implement tabular Q-learning from scratch for a gridworld environment.
- Compare Q-learning (off-policy) and SARSA (on-policy) in terms of learned policies.
- Analyze exploration strategies: ϵ -greedy versus upper confidence bound (UCB).

Motivation

An agent navigates a maze, choosing directions and receiving rewards, with no prior knowledge of the map. This is the **reinforcement learning** (RL) problem [sutton2018]. Unlike supervised learning (correct answers provided) or game theory’s analytical equilibria, RL learns through trial-and-error interaction. The framework is foundational for multi-agent learning: before understanding how two Q-learners interact in a matrix game (sec-q-learning), we must understand single-agent Q-learning.

Theory

Markov decision processes

A **Markov decision process** (MDP) is a tuple $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$ where:

- $\mathcal{S}$ is a finite set of states,
- $\mathcal{A}$ is a finite set of actions,
- $P(s' | s, a)$ is the transition probability,
- $R(s, a, s')$ is the reward function,
- $\gamma \in [0, 1)$ is the discount factor.

The agent’s goal is to find a policy $\pi(a | s)$ that maximizes expected cumulative discounted reward:

$$V^\pi(s) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t R_t \mid S_0 = s \right] \text{ (\#eq : value - function)} \quad (1)$$

The Bellman equation

The optimal value function $V^*(s)$ satisfies the **Bellman optimality equation**:

$$V^*(s) = \max_{a \in \mathcal{A}} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma V^*(s')] \text{ (\#eq : bellman)} \quad (2)$$

Equivalently, the optimal action-value function $Q^*(s, a)$ satisfies:

$$Q^*(s, a) = \sum_{s'} P(s' | s, a) \left[R(s, a, s') + \gamma \max_{a'} Q^*(s', a') \right] \text{ (\#eq : bellman - q)} \quad (3)$$

Tabular Q-learning

Q-learning [Watkins1992] is a model-free, off-policy algorithm. After taking action a in state s , observing reward r and next state s' :

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] \quad (\#eq : q - update) \quad (4)$$

The max operator makes it off-policy: Q-learning updates toward the best possible next action regardless of what was actually taken. Under standard conditions, $Q \rightarrow Q^*$.

SARSA

SARSA (State-Action-Reward-State-Action) is an on-policy alternative that bootstraps from the action a' actually chosen, rather than the greedy max:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)] \quad (\#eq : sarsa) \quad (5)$$

SARSA learns the value of the policy it follows (including exploration), making it more conservative near dangerous states.

Exploration vs exploitation

The agent must balance **exploitation** (best-known action) with **exploration** (trying alternatives). Two standard strategies: **ϵ -greedy** (random action with probability ϵ , greedy otherwise) and **UCB** (choose $a = \arg \max_a [Q(s, a) + c\sqrt{\ln t / N(s, a)}]$, exploring in proportion to uncertainty).

Implementation in R

Gridworld environment

```
create_gridworld <- function(rows = 5, cols = 5, obstacles = NULL,
                             goal = NULL, goal_reward = 10, step_cost = -0.1) {
  if (is.null(goal)) goal <- c(rows, cols)
  # State encoding: (row, col) -> integer id
  state_id <- function(r, c) (r - 1) * cols + c
  id_to_rc <- function(id) c((id - 1) %/% cols + 1, (id - 1) %% cols + 1)

  is_obstacle <- function(r, c) {
    if (is.null(obstacles)) return(FALSE)
    any(obstacles[, 1] == r & obstacles[, 2] == c)
  }

  # Actions: 1=up, 2=right, 3=down, 4=left
  action_delta <- list(c(-1, 0), c(0, 1), c(1, 0), c(0, -1))
  action_names <- c("Up", "Right", "Down", "Left")

  step <- function(state, action) {
    rc <- id_to_rc(state)
    delta <- action_delta[[action]]
```

```

nr <- rc[1] + delta[1]
nc <- rc[2] + delta[2]

# Stay in place if hitting wall or obstacle
if (nr < 1 || nr > rows || nc < 1 || nc > cols || is_obstacle(nr, nc)) {
  return(list(next_state = state, reward = step_cost, done = FALSE))
}

next_state <- state_id(nr, nc)
done <- (nr == goal[1] && nc == goal[2])
reward <- if (done) goal_reward else step_cost

list(next_state = next_state, reward = reward, done = done)
}

list(
  rows = rows, cols = cols, n_states = rows * cols, n_actions = 4,
  start = state_id(1, 1), goal = state_id(goal[1], goal[2]),
  step = step, state_id = state_id, id_to_rc = id_to_rc,
  action_names = action_names, obstacles = obstacles
)
}

```

Tabular Q-learning agent

```

q_learning_gridworld <- function(env, episodes = 1000,
                                alpha = 0.1, gamma = 0.95,
                                epsilon = 0.1, max_steps = 200) {
  Q <- matrix(0, nrow = env$n_states, ncol = env$n_actions)
  episode_rewards <- numeric(episodes)

  for (ep in seq_len(episodes)) {
    state <- env$start
    total_reward <- 0

    for (step in seq_len(max_steps)) {
      # Epsilon-greedy action selection
      if (runif(1) < epsilon) {
        action <- sample(env$n_actions, 1)
      } else {
        ties <- which(Q[state, ] == max(Q[state, ]))
        action <- sample(ties, 1)
      }

      result <- env$step(state, action)
      # Q-learning update (off-policy: uses max over next actions)
      best_next <- max(Q[result$next_state, ])
      Q[state, action] <- Q[state, action] +
        alpha * (result$reward + gamma * best_next - Q[state, action])

      total_reward <- total_reward + result$reward
      state <- result$next_state
    }
  }
}

```

```

    if (result$done) break
  }
  episode_rewards[ep] <- total_reward
}

list(Q = Q, rewards = episode_rewards)
}

```

SARSA agent

```

sarsa_gridworld <- function(env, episodes = 1000,
                             alpha = 0.1, gamma = 0.95,
                             epsilon = 0.1, max_steps = 200) {
  Q <- matrix(0, nrow = env$n_states, ncol = env$n_actions)
  episode_rewards <- numeric(episodes)

  choose_action <- function(state) {
    if (runif(1) < epsilon) {
      sample(env$n_actions, 1)
    } else {
      ties <- which(Q[state, ] == max(Q[state, ]))
      sample(ties, 1)
    }
  }

  for (ep in seq_len(episodes)) {
    state <- env$start
    action <- choose_action(state)
    total_reward <- 0

    for (step in seq_len(max_steps)) {
      result <- env$step(state, action)
      next_action <- choose_action(result$next_state)

      # SARSA update (on-policy: uses actual next action)
      Q[state, action] <- Q[state, action] +
        alpha * (result$reward + gamma * Q[result$next_state, next_action] -
                  Q[state, action])

      total_reward <- total_reward + result$reward
      state <- result$next_state
      action <- next_action
      if (result$done) break
    }
    episode_rewards[ep] <- total_reward
  }

  list(Q = Q, rewards = episode_rewards)
}

```

Training on the gridworld

```
obstacles <- matrix(c(2, 2, 3, 4, 4, 2, 2, 4), ncol = 2, byrow = TRUE)

env <- create_gridworld(rows = 5, cols = 5, obstacles = obstacles)

set.seed(42)
ql_result <- q_learning_gridworld(env, episodes = 2000,
                                  alpha = 0.15, gamma = 0.95, epsilon = 0.1)
set.seed(42)
sarsa_result <- sarsa_gridworld(env, episodes = 2000,
                                alpha = 0.15, gamma = 0.95, epsilon = 0.1)

# Extract max Q-value per state as value function
V <- apply(ql_result$Q, 1, max)

grid_df <- tibble(
  state = seq_len(env$n_states),
  row = (state - 1) %/% env$cols + 1,
  col = (state - 1) %% env$cols + 1,
  value = V
)

# Mark obstacles
obs_df <- tibble(
  row = obstacles[, 1],
  col = obstacles[, 2]
)

p_heatmap <- ggplot(grid_df, aes(x = col, y = row, fill = value)) +
  geom_tile(colour = "white", linewidth = 0.5) +
  geom_tile(data = obs_df, aes(x = col, y = row),
            fill = "grey20", inherit.aes = FALSE) +
  geom_text(aes(label = round(value, 1)), colour = "black", size = 3.5) +
  annotate("text", x = 1, y = 1, label = "S", fontface = "bold", size = 5) +
  annotate("text", x = 5, y = 5, label = "G", fontface = "bold", size = 5) +
  scale_fill_gradient(low = "white", high = okabe_ito[1], name = "V(s)") +
  scale_y_reverse() +
  coord_equal() +
  theme_publication() +
  theme(panel.grid = element_blank()) +
  labs(title = "Q-Learning: Learned State Values",
       x = "Column", y = "Row")

p_heatmap

save_pub_fig(p_heatmap, "q-learning-gridworld-heatmap", width = 6, height = 5)
```

Q-Learning: Learned State Values




Figure 1: Learned state values (max Q per state) in the 5x5 gridworld. The agent starts at (1,1) and must reach the goal at (5,5). Dark cells are obstacles. Higher values (orange) indicate states closer to the goal along the optimal path.

Exploration comparison: epsilon-greedy vs UCB

```
run_bandit <- function(strategy, n_arms = 10, n_steps = 2000,
                        epsilon = 0.1, ucb_c = 2) {
  true_values <- rnorm(n_arms); Q <- rep(0, n_arms); N <- rep(0, n_arms)
  cum_regret <- numeric(n_steps); optimal <- max(true_values); tot <- 0
  for (t in seq_len(n_steps)) {
    action <- if (strategy == "epsilon_greedy") {
      if (runif(1) < epsilon) sample(n_arms, 1) else which.max(Q)
    } else {
      if (any(N == 0)) which(N == 0)[1]
      else which.max(Q + ucb_c * sqrt(log(t) / N))
    }
    reward <- rnorm(1, mean = true_values[action])
    N[action] <- N[action] + 1
    Q[action] <- Q[action] + (reward - Q[action]) / N[action]
    tot <- tot + (optimal - true_values[action]); cum_regret[t] <- tot
  }
  cum_regret
}

set.seed(42); n_runs <- 50; n_steps <- 2000
eg_regret <- matrix(0, n_runs, n_steps); ucb_regret <- matrix(0, n_runs, n_steps)
for (i in seq_len(n_runs)) {
  eg_regret[i, ] <- run_bandit("epsilon_greedy", n_steps = n_steps)
  ucb_regret[i, ] <- run_bandit("ucb", n_steps = n_steps)
}

regret_df <- tibble(
  step = rep(1:n_steps, 2),
  regret = c(colMeans(eg_regret), colMeans(ucb_regret)),
  strategy = rep(c("Epsilon-greedy", "UCB"), each = n_steps)
)

p_explore <- ggplot(regret_df, aes(x = step, y = regret, colour = strategy)) +
  geom_line(linewidth = 0.8) +
  scale_colour_manual(values = c("Epsilon-greedy" = okabe_ito[1],
                                "UCB" = okabe_ito[2])) +
  theme_publication() +
  labs(title = "Exploration: Epsilon-Greedy vs UCB",
       x = "Step", y = "Cumulative regret", colour = "Strategy")
p_explore

save_pub_fig(p_explore, "exploration-eg-vs-ucb", width = 7, height = 5)
```

Worked example

We trace the learning process on the 5x5 gridworld with obstacles at (2,2), (3,4), (4,2), and (2,4). The agent starts at (1,1) and must reach the goal at (5,5).




Figure 2: Cumulative regret comparison between epsilon-greedy and UCB on a 10-armed bandit. UCB explores proportionally to uncertainty, yielding lower regret.

```
cat("Q-Learning - Mean reward (last 200 episodes):",
    round(mean(tail(ql_result$rewards, 200)), 2), "\n")
```

```
#> Q-Learning - Mean reward (last 200 episodes): 6.26
```

```
cat("SARSA - Mean reward (last 200 episodes):",
    round(mean(tail(sarsa_result$rewards, 200)), 2), "\n")
```

```
#> SARSA - Mean reward (last 200 episodes): 6.28
```

```
# Extract and display the learned policy
policy <- apply(ql_result$Q, 1, which.max)
arrows <- c("Up", "Rt", "Dn", "Lt")
obs_states <- env$state_id(obstacles[,1], obstacles[,2])

policy_matrix <- matrix("", nrow = env$rows, ncol = env$cols)
for (s in seq_len(env$n_states)) {
  r <- (s - 1) %% env$cols + 1
  cc <- (s - 1) %% env$cols + 1
  if (s %in% obs_states) {
    policy_matrix[r, cc] <- " X "
  } else if (s == env$goal) {
    policy_matrix[r, cc] <- " G "
  } else {
    policy_matrix[r, cc] <- arrows[policy[s]]
  }
}
```



```

    }
}
cat("\nLearned policy (Q-learning):\n")

```

```

#>
#> Learned policy (Q-learning):

```

```

print(policy_matrix, quote = FALSE)

```

```

#>      [,1] [,2] [,3] [,4] [,5]
#> [1,] Rt  Rt  Rt  Rt  Dn
#> [2,] Up   X   Up   X   Dn
#> [3,] Dn   Dn  Up   X   Dn
#> [4,] Lt   X   Rt  Rt  Dn
#> [5,] Rt  Rt  Rt  Rt  G

```

Q-learning finds the optimal route because it evaluates the greedy policy (via max). SARSA gives obstacles a wider berth because its on-policy updates incorporate the risk of accidentally stepping into a bad state during exploration.

Extensions

- **Function approximation:** For large state spaces, Q-values can be approximated with linear models or neural networks (DQN). Neural network packages are beyond our R toolkit, but the conceptual framework is identical.
- **Policy gradient methods:** REINFORCE and actor-critic methods directly optimize policy parameters rather than learning value functions.
- **Multi-agent extensions:** @ref(sec-q-learning) explores what happens when two Q-learners face each other in matrix games — single-agent convergence guarantees break down.
- **Planning:** Model-based RL (Dyna-Q) combines learned models with Q-learning for faster convergence. See @sutton2018 (Chapters 5-8).

Exercises

1. **Cliff walking.** Create a 4x12 gridworld where the bottom row (except the start and goal) is a “cliff” with reward -100 . Compare Q-learning and SARSA policies. Which algorithm learns the safer path along the top?
2. **Decaying epsilon.** Implement $\varepsilon_t = 1/\sqrt{t}$ decaying exploration for the 5x5 gridworld. Plot the learning curve alongside fixed $\varepsilon = 0.1$ and compare convergence speed and final performance.
3. **Discount factor analysis.** Train Q-learning on the 5x5 gridworld with $\gamma \in \{0.5, 0.9, 0.99\}$. Visualize the learned value functions as heatmaps. How does the discount factor affect the value gradient from start to goal?

Solutions appear in @ref(sec-solutions).