

Multi-Agent RL

Independent learners, joint-action learners, minimax-Q for zero-sum games, and the conceptual foundations of MADDPG and PSRO.

Learning objectives

- Distinguish independent learners from joint-action learners and explain why joint-action learning addresses non-stationarity.
- Implement minimax-Q learning for two-player zero-sum games in base R and show convergence to the minimax strategy.
- Describe the centralized-training, decentralized-execution paradigm underlying MADDPG.
- Explain how Policy Space Response Oracles (PSRO) combine reinforcement learning with game-theoretic solution concepts.

Motivation

Chapter @ref(sec-q-learning) showed that independent Q-learners converge in coordination games but cycle in zero-sum games like Matching Pennies. This failure is not a bug in the algorithm — it reflects a fundamental limitation. Each agent treats the other as part of a stationary environment, but both are simultaneously adapting. The environment is anything but stationary.

Multi-agent reinforcement learning (MARL) addresses this challenge through algorithms that explicitly account for the presence of other learners. The simplest upgrade is **joint-action learning**, where each agent observes the opponent’s actions and maintains Q-values over joint action profiles. More sophisticated approaches include **minimax-Q** [littman1994], which finds optimal strategies in zero-sum games, and **MADDPG** [lowe2017], which scales to continuous action spaces. At the frontier, **PSRO** [lanctot2017] bridges MARL with classical game-theoretic solution concepts.

Understanding these approaches is essential for anyone working on competitive AI systems, from game-playing agents to adversarial robustness in machine learning.

Theory

Independent vs. joint-action learners

An **independent learner** (IL) maintains Q-values $Q_i(a_i)$ over its own actions only, as in @ref(sec-q-learning). It ignores the opponent entirely.

A **joint-action learner** (JAL) observes the opponent’s action after each round and maintains Q-values over the joint action profile: $Q_i(a_i, a_{-i})$. The agent also tracks an empirical model of the opponent’s action frequencies $\hat{\pi}_{-i}(a_{-i})$ and selects actions that maximize expected payoff under that model:

$$a_i^* = \arg \max_{a_i} \sum_{a_{-i}} \hat{\pi}_{-i}(a_{-i}) Q_i(a_i, a_{-i}) (\#eq : jal - action) \quad (1)$$

JAL addresses non-stationarity by adapting its strategy to the opponent’s observed behavior. However, it still assumes the opponent’s policy is stationary — an assumption that breaks when both agents are learning.

Minimax-Q learning

For two-player zero-sum games, @littman1994 proposed **minimax-Q**, which replaces the max operator in Q-learning with a minimax calculation. Instead of learning a deterministic best action, the agent learns a **mixed strategy** π_i that maximizes its worst-case expected payoff:

$$V_i = \max_{\pi_i} \min_{a_{-i}} \sum_{a_i} \pi_i(a_i) Q_i(a_i, a_{-i}) (\#eq : minimax - value) \quad (2)$$

The Q-value update becomes:

$$Q_i(a_i, a_{-i}) \leftarrow Q_i(a_i, a_{-i}) + \alpha [r_i + \gamma V_i - Q_i(a_i, a_{-i})] (\#eq : minimax - q - update) \quad (3)$$

In the stateless (matrix game) setting, $\gamma = 0$ and the update simplifies to:

$$Q_i(a_i, a_{-i}) \leftarrow Q_i(a_i, a_{-i}) + \alpha [r_i - Q_i(a_i, a_{-i})] (\#eq : minimax - q - stateless) \quad (4)$$

The key insight is that Q-values converge to the game’s payoff matrix, and the minimax strategy over those Q-values converges to the Nash equilibrium mixed strategy.

MADDPG: centralized training, decentralized execution

Multi-Agent Deep Deterministic Policy Gradient [lowe2017] extends actor-critic methods to multi-agent settings. The core idea is **centralized training with decentralized execution**:

- **During training**, each agent’s *critic* has access to the actions and observations of all agents. This makes the environment stationary from the critic’s perspective — it sees the full joint state.
- **During execution**, each agent’s *actor* uses only its own observations to select actions. No communication is needed at deployment time.

This paradigm elegantly sidesteps the non-stationarity problem during learning while maintaining practical deployability. MADDPG requires neural network function approximation (beyond our available packages), but the architectural principle — train with global information, execute with local information — is broadly applicable.

Policy Space Response Oracles (PSRO)

PSRO [lanctot2017] takes a fundamentally different approach. Instead of training a single policy per agent, it maintains a **population of policies** and iteratively expands it:

1. Start with an initial policy for each agent.
2. Compute a **meta-game** payoff matrix: the expected payoff when each agent plays each policy in its population.
3. Solve the meta-game for a **meta-strategy** (e.g., a Nash equilibrium over the policy populations).
4. Train a new **best response** policy for each agent against the opponent’s current meta-strategy.
5. Add the new policies to the populations and repeat.

PSRO unifies several earlier approaches. When the meta-strategy solver is a Nash equilibrium computation and the best-response oracle is an RL algorithm, PSRO recovers the **double oracle** method from classical game theory. The framework has been applied to large-scale games including StarCraft [vinyals2019].

Implementation in R

Joint-action learner

```
jal_agent_create <- function(n_own, n_opp, alpha = 0.1, epsilon = 0.1) {  
  Q <- matrix(0, nrow = n_own, ncol = n_opp)  
  opp_counts <- rep(1, n_opp) # Laplace smoothing  
  
  list(  
    choose = function() {  
      opp_freq <- opp_counts / sum(opp_counts)  
      if (runif(1) < epsilon) {  
        sample(n_own, 1)  
      } else {  
        ev <- Q %*% opp_freq  
        ties <- which(ev == max(ev))  
        sample(ties, 1)  
      }  
    },  
    update = function(own_action, opp_action, reward) {  
      Q[own_action, opp_action] <-<- Q[own_action, opp_action] +  
        alpha * (reward - Q[own_action, opp_action])  
      opp_counts[opp_action] <-<- opp_counts[opp_action] + 1  
    },  
    get_Q = function() Q,  
    get_opp_model = function() opp_counts / sum(opp_counts)  
  )  
}
```

Minimax-Q solver

For a 2×2 zero-sum game, the minimax mixed strategy has a closed-form solution. The row player's optimal probability on action 1 is:

```
solve_minimax_2x2 <- function(Q) {  
  # Q is a 2x2 matrix of row player's Q-values  
  # Returns optimal mixed strategy (probability on action 1)  
  denom <- Q[1, 1] - Q[1, 2] - Q[2, 1] + Q[2, 2]  
  if (abs(denom) < 1e-10) return(0.5)  
  p <- (Q[2, 2] - Q[2, 1]) / denom  
  p <- max(0, min(1, p)) # clamp to [0, 1]  
  p  
}
```

Minimax-Q agent

```
minimax_q_create <- function(n_own, n_opp, alpha = 0.1) {  
  Q <- matrix(0, nrow = n_own, ncol = n_opp)  
  
  list(  

```

```

choose = function() {
  p <- solve_minimax_2x2(Q)
  if (runif(1) < p) 1L else 2L
},
update = function(own_action, opp_action, reward) {
  Q[own_action, opp_action] <- Q[own_action, opp_action] +
    alpha * (reward - Q[own_action, opp_action])
},
get_Q = function() Q,
get_strategy = function() {
  p <- solve_minimax_2x2(Q)
  c(p, 1 - p)
}
)
}

```

Simulation: independent vs. joint-action learners

```

il_agent_create <- function(n_actions, alpha = 0.1, epsilon = 0.1) {
  env <- new.env(parent = emptyenv())
  env$Q <- rep(0, n_actions)
  list(
    choose = function() {
      if (runif(1) < epsilon) sample(n_actions, 1)
      else { ties <- which(env$Q == max(env$Q)); sample(ties, 1) }
    },
    update = function(a, opp_a, r) {
      env$Q[a] <- env$Q[a] + alpha * (r - env$Q[a])
    }
  )
}

run_il_vs_jal <- function(A, B, episodes = 5000, alpha = 0.1,
  epsilon = 0.1, type = "IL") {
  n1 <- nrow(A); n2 <- ncol(A)
  if (type == "IL") {
    agent1 <- il_agent_create(n1, alpha, epsilon)
    agent2 <- il_agent_create(n2, alpha, epsilon)
  } else {
    agent1 <- jal_agent_create(n1, n2, alpha, epsilon)
    agent2 <- jal_agent_create(n2, n1, alpha, epsilon)
  }
  rewards <- numeric(episodes)
  for (ep in seq_len(episodes)) {
    a1 <- agent1$choose(); a2 <- agent2$choose()
    r1 <- A[a1, a2]; r2 <- B[a1, a2]
    agent1$update(a1, a2, r1); agent2$update(a2, a1, r2)
    rewards[ep] <- r1 + r2
  }
  rewards
}

```

```

# Coordination game: both players want to match
A_coord <- matrix(c(2, 0, 0, 1), nrow = 2, byrow = TRUE)
B_coord <- A_coord

set.seed(42)
n_runs <- 20
episodes <- 3000
window <- 100

il_runs <- replicate(n_runs, run_il_vs_jal(A_coord, B_coord,
                                           episodes = episodes, type = "IL"))
jal_runs <- replicate(n_runs, run_il_vs_jal(A_coord, B_coord,
                                           episodes = episodes, type = "JAL"))

smooth_rewards <- function(runs, window) {
  apply(runs, 2, function(r) {
    stats::filter(r, rep(1 / window, window), sides = 1)
  }) |>
  rowMeans(na.rm = TRUE)
}

df_curves <- tibble(
  episode = rep(seq_len(episodes), 2),
  reward = c(smooth_rewards(il_runs, window),
             smooth_rewards(jal_runs, window)),
  learner = rep(c("Independent", "Joint-Action"), each = episodes)
) |>
  filter(!is.na(reward))

p_curves <- ggplot(df_curves, aes(x = episode, y = reward, colour = learner)) +
  geom_line(linewidth = 0.8) +
  geom_hline(yintercept = 4, linetype = "dashed", colour = "grey50") +
  annotate("text", x = 2800, y = 3.85, label = "Optimal (both play A)",
          size = 3, colour = "grey40") +
  scale_colour_manual(values = c("Independent" = okabe_ito[1],
                                "Joint-Action" = okabe_ito[2])) +
  theme_publication() +
  labs(title = "Independent vs. Joint-Action Learners",
       x = "Episode", y = "Joint Payoff (smoothed)",
       colour = "Learner Type")

p_curves

save_pub_fig(p_curves, "marl-learning-curves", width = 7, height = 5)

```

Win-rate matrix across strategies

```

# Matching Pennies tournament: IL vs JAL vs Minimax-Q
play_match <- function(create_p1, create_p2, A, B, episodes = 5000) {
  p1 <- create_p1()

```




Figure 1: Learning curves in a coordination game (joint payoff, smoothed over 100 episodes, averaged over 20 runs). Joint-action learners coordinate faster and more reliably than independent learners because they model the opponent's action frequencies.

```

p2 <- create_p2()
wins <- 0
for (ep in seq_len(episodes)) {
  a1 <- p1$choose()
  a2 <- p2$choose()
  r1 <- A[a1, a2]
  r2 <- B[a1, a2]
  p1$update(a1, a2, r1)
  p2$update(a2, a1, r2)
  if (r1 > 0) wins <- wins + 1
}
wins / episodes
}

A_mp <- matrix(c(1, -1, -1, 1), nrow = 2, byrow = TRUE)
B_mp <- -A_mp

make_il <- function() il_agent_create(2, 0.1, 0.1)
make_jal <- function() jal_agent_create(2, 2, 0.1, 0.1)
make_minimax <- function() {
  a <- minimax_q_create(2, 2, 0.1)
  list(choose = a$choose, update = a$update)
}

set.seed(123)

```

```

strategies <- list(IL = make_il, JAL = make_jal, Minimax = make_minimax)
n_strats <- length(strategies)
win_matrix <- matrix(0, n_strats, n_strats,
                     dimnames = list(names(strategies), names(strategies)))

n_matches <- 30
for (i in seq_len(n_strats)) {
  for (j in seq_len(n_strats)) {
    rates <- replicate(n_matches,
                       play_match(strategies[[i]], strategies[[j]], A_mp, B_mp, episodes = 2000))
    win_matrix[i, j] <- mean(rates)
  }
}

```

```

df_heat <- expand_grid(
  row_strategy = factor(rownames(win_matrix), levels = rownames(win_matrix)),
  col_strategy = factor(colnames(win_matrix), levels = colnames(win_matrix))
) |>
  mutate(win_rate = as.vector(win_matrix))

p_heat <- ggplot(df_heat, aes(x = col_strategy, y = row_strategy,
                             fill = win_rate)) +
  geom_tile(colour = "white", linewidth = 1.5) +
  geom_text(aes(label = sprintf("%.1f%%", win_rate * 100)), size = 4.5) +
  scale_fill_gradient2(low = okabe_ito[2], mid = "white", high = okabe_ito[1],
                      midpoint = 0.5, limits = c(0.3, 0.7),
                      labels = scales::percent) +
  theme_publication() +
  theme(panel.grid = element_blank()) +
  labs(title = "Win Rate Matrix: Matching Pennies",
       x = "Column Player Strategy", y = "Row Player Strategy",
       fill = "Row Player Win Rate")

p_heat

```

```

save_pub_fig(p_heat, "marl-win-rate-matrix", width = 6, height = 5)

```

Worked example

We trace minimax-Q learning in Matching Pennies, showing convergence to the minimax strategy of playing each action with probability 0.5.

```

set.seed(42)
agent <- minimax_q_create(2, 2, alpha = 0.05)
episodes <- 5000

strategy_history <- matrix(NA, nrow = episodes, ncol = 2)
q_history <- matrix(NA, nrow = episodes, ncol = 4)

for (ep in seq_len(episodes)) {
  a1 <- agent$choose()

```




Figure 2: Win-rate matrix in Matching Pennies (row player’s win rate). Minimax-Q achieves near 50% against all opponents — the theoretically optimal rate in a zero-sum game — while independent and joint-action learners show exploitable biases.


```

a2 <- sample(2, 1) # opponent plays uniformly
r1 <- A_mp[a1, a2]
agent$update(a1, a2, r1)
strategy_history[ep, ] <- agent$get_strategy()
Q <- agent$get_Q()
q_history[ep, ] <- as.vector(Q)
}

cat("Minimax-Q convergence in Matching Pennies:\n\n")

```

```
#> Minimax-Q convergence in Matching Pennies:
```

```

checkpoints <- c(100, 500, 1000, 2000, 5000)
for (ep in checkpoints) {
  s <- strategy_history[ep, ]
  cat(sprintf(" Episode %4d: P(Heads) = %.3f, P(Tails) = %.3f\n",
              ep, s[1], s[2]))
}

```

```

#> Episode 100: P(Heads) = 0.516, P(Tails) = 0.484
#> Episode 500: P(Heads) = 0.500, P(Tails) = 0.500
#> Episode 1000: P(Heads) = 0.500, P(Tails) = 0.500
#> Episode 2000: P(Heads) = 0.500, P(Tails) = 0.500
#> Episode 5000: P(Heads) = 0.500, P(Tails) = 0.500

```

```
cat(sprintf("\nFinal Q-values:\n"))
```

```

#>
#> Final Q-values:

```

```

Q_final <- agent$get_Q()
rownames(Q_final) <- c("H", "T")
colnames(Q_final) <- c("H", "T")
print(round(Q_final, 3))

```

```

#>   H  T
#> H  1 -1
#> T -1  1

```

```
cat(sprintf("\nThe Nash equilibrium is P(Heads) = 0.500."))
```

```

#>
#> The Nash equilibrium is P(Heads) = 0.500.

```

```

cat(sprintf("\nFinal strategy: P(Heads) = %.3f --- convergence achieved.\n",
            strategy_history[episodes, 1]))

```

```

#>
#> Final strategy: P(Heads) = 0.500 --- convergence achieved.

```

The minimax-Q agent converges to a strategy near $(0.5, 0.5)$. Unlike independent Q-learners (which cycle endlessly in this game, as shown in @ref(sec-q-learning)), the minimax formulation explicitly optimizes for worst-case performance. The Q-values converge to approximate the true payoff matrix, and the minimax calculation over those Q-values yields the Nash equilibrium mixed strategy.

The key advantage is robustness: the minimax strategy guarantees an expected payoff of zero regardless of the opponent’s play. No deterministic strategy can make this guarantee.

Extensions

- **Win-or-Learn-Fast (WoLF-PHC):** Varies the learning rate depending on whether the agent is currently winning or losing relative to its equilibrium payoff. This achieves convergence in a wider class of games than standard Q-learning.
- **Nash-Q:** Extends minimax-Q to general-sum games by computing a Nash equilibrium at each update step. Computationally expensive but theoretically more general.
- **Opponent modelling:** Rather than assuming worst-case play, agents can explicitly model the opponent’s learning dynamics and best-respond to predicted future behavior.
- **Self-play** (@ref(sec-self-play)) combines RL with self-play to learn without any opponent model, scaling to games like Go and chess.
- **Counterfactual regret minimization** (@ref(sec-cfr)) provides an alternative to RL-based approaches for imperfect-information games.

For a comprehensive survey of MARL, see @shoham2009. The PSRO framework is detailed in @lanctot2017.

Exercises

1. **Joint-action learner in Prisoner’s Dilemma.** Implement two joint-action learners playing the Prisoner’s Dilemma ($T = 5, R = 3, P = 1, S = 0$) for 5000 episodes. Does the JAL converge to mutual defection? Compare the convergence speed to independent learners from @ref(sec-q-learning).
2. **Minimax-Q with decaying learning rate.** Modify the minimax-Q agent to use a decaying learning rate $\alpha_t = 1/(1 + 0.001 \cdot t)$. Run the Matching Pennies experiment again. Does the decaying rate improve convergence stability? Plot the strategy trajectory $P(\text{Heads})$ over episodes for both constant and decaying rates.
3. **PSRO by hand.** Consider Rock-Paper-Scissors. Start with a population of two policies for each player: “always Rock” and “always Paper.” Compute the 2×2 meta-game payoff matrix. Find the Nash equilibrium of this meta-game. What best-response policy should be added next? Repeat for one more iteration and discuss whether the meta-strategy converges toward the full-game Nash equilibrium of $(1/3, 1/3, 1/3)$.

Solutions appear in @ref(sec-solutions).