

Mechanism Design

Inverse game theory: the revelation principle, incentive compatibility, the VCG mechanism, and implementation in R for public goods provision.

Learning objectives

- Explain mechanism design as the “inverse” problem of game theory: designing rules to achieve desired outcomes.
- State the revelation principle and its role in simplifying mechanism design.
- Define incentive compatibility and individual rationality.
- Implement the VCG mechanism for a public goods provision problem in R.

Motivation

In standard game theory, we take the rules of the game as given and predict outcomes. Mechanism design flips this: given a desired outcome, what rules should we impose? This “reverse engineering” perspective is central to real-world institutional design — from spectrum auctions and organ allocation systems to voting rules and matching markets.

Consider a city council deciding whether to build a public park. Each citizen has a private valuation for the park, which they may misrepresent if doing so serves their interest. Can we design a mechanism — a set of rules for eliciting preferences and making decisions — that induces truthful reporting and achieves the socially optimal outcome? The Vickrey-Clarke-Groves (VCG) mechanism shows that the answer is yes, under certain conditions [Nisan2007, Chapter 9].

Theory

The mechanism design problem

A **mechanism** specifies (i) a message space for each agent and (ii) an outcome function mapping messages to decisions and transfers. A mechanism **implements** a social choice function if, in equilibrium, the mechanism’s outcome matches what the social choice function prescribes.

The revelation principle

Theorem: Revelation Principle

For any mechanism in which agents play a Bayesian Nash equilibrium, there exists a **direct** mechanism — one where agents simply report their types — that achieves the same outcome and where truthful reporting is an equilibrium.

The revelation principle dramatically simplifies the designer’s problem: instead of searching over all possible mechanisms, we can restrict attention to direct, truthful mechanisms without loss of generality.

Incentive compatibility and individual rationality

A mechanism is **incentive compatible (IC)** if truthful reporting is an equilibrium strategy for every agent. It is **individually rational (IR)** if every agent’s expected payoff from participating is at least as large as their outside option (typically zero).

These two constraints define the boundary of what is achievable. A mechanism that violates IC will produce distorted reports; one that violates IR will see agents opt out.

The VCG mechanism

The Vickrey-Clarke-Groves mechanism works as follows for a public decision:

1. Each agent i reports a valuation $\hat{v}_i$ for each possible outcome.
2. The mechanism chooses the outcome k^* that maximizes total reported welfare: $k^* = \arg \max_k \sum_i \hat{v}_i(k)$.
3. Each agent pays a **Clarke tax**: the externality they impose on others.

$$t_i = \max_k \sum_{j \neq i} \hat{v}_j(k) - \sum_{j \neq i} \hat{v}_j(k^*) (\#eq : vcg - tax) \quad (1)$$

The Clarke tax equals the difference between what the other agents *could* have achieved (without i) and what they *do* achieve (with i 's presence influencing the decision). Under VCG, truthful reporting is a dominant strategy [nisan2007, Chapter 9].

Implementation in R

VCG mechanism for public goods

```
vcg_mechanism <- function(valuations) {  
  # valuations: matrix where rows = agents, cols = outcomes  
  # Returns: chosen outcome, transfers, and payoffs  
  n_agents <- nrow(valuations)  
  n_outcomes <- ncol(valuations)  
  
  # Social welfare for each outcome  
  social_welfare <- colSums(valuations)  
  chosen <- which.max(social_welfare)  
  
  # Clarke tax for each agent  
  transfers <- numeric(n_agents)  
  for (i in seq_len(n_agents)) {  
    # Best outcome for others without agent i's influence  
    welfare_without_i <- colSums(valuations[-i, , drop = FALSE])  
    best_without_i <- max(welfare_without_i)  
    # What others get under the chosen outcome  
    others_at_chosen <- sum(valuations[-i, chosen])  
    transfers[i] <- best_without_i - others_at_chosen  
  }  
  
  # Agent payoff = valuation of chosen outcome - tax  
  payoffs <- valuations[, chosen] - transfers  
  
  list(  
    chosen_outcome = chosen,  
    social_welfare = social_welfare[chosen],  
    transfers = transfers,  
    payoffs = payoffs,  
    is_truthful_dominant = TRUE # by VCG theory  
  )  
}
```

```

# Majority voting mechanism (for comparison)
majority_voting <- function(valuations) {
  n_agents <- nrow(valuations)
  n_outcomes <- ncol(valuations)

  # Each agent votes for their most-preferred outcome
  votes <- apply(valuations, 1, which.max)
  vote_counts <- tabulate(votes, nbins = n_outcomes)
  chosen <- which.max(vote_counts)

  list(
    chosen_outcome = chosen,
    social_welfare = sum(valuations[, chosen]),
    transfers = rep(0, n_agents),
    payoffs = valuations[, chosen]
  )
}

```

Comparing mechanisms across many instances

```

set.seed(42)
n_simulations <- 500
n_agents <- 3
n_outcomes <- 2 # Build park (1) vs. don't build (2)

results <- map_dfr(seq_len(n_simulations), function(sim_id) {
  # Random valuations: agents value "build" differently, "don't build" = 0
  v_build <- runif(n_agents, min = -2, max = 5)
  valuations <- cbind(v_build, rep(0, n_agents))

  vcg_result <- vcg_mechanism(valuations)
  mv_result <- majority_voting(valuations)

  # Optimal social welfare (first-best)
  optimal_welfare <- max(colSums(valuations))

  tibble(
    sim_id = sim_id,
    mechanism = c("VCG", "Majority voting", "First-best"),
    welfare = c(vcg_result$social_welfare,
                mv_result$social_welfare,
                optimal_welfare)
  )
})

```

Figure 1: Social welfare comparison

```

welfare_summary <- results |>
  group_by(mechanism) |>

```

```

summarise(
  mean_welfare = mean(welfare),
  se_welfare = sd(welfare) / sqrt(n()),
  .groups = "drop"
) |>
mutate(mechanism = factor(mechanism,
                          levels = c("First-best", "VCG", "Majority voting")))

p_welfare <- ggplot(welfare_summary,
                   aes(x = mechanism, y = mean_welfare, fill = mechanism)) +
  geom_col(alpha = 0.85, width = 0.6) +
  geom_errorbar(aes(ymin = mean_welfare - 1.96 * se_welfare,
                    ymax = mean_welfare + 1.96 * se_welfare),
               width = 0.15) +
  geom_text(aes(label = sprintf("%.2f", mean_welfare)),
            vjust = -0.8, size = 3.5) +
  scale_fill_manual(values = okabe_ito[c(3, 1, 2)]) +
  scale_y_continuous(name = "Mean social welfare",
                     expand = expansion(mult = c(0, 0.15))) +
  labs(x = "Mechanism",
       title = "Social Welfare Across Mechanisms",
       subtitle = glue("{n_simulations} simulations, {n_agents} agents, ",
                       "public goods decision")) +
  theme_publication() +
  theme(legend.position = "none")

p_welfare

save_pub_fig(p_welfare, "mechanism-welfare-comparison", width = 6, height = 5)

```

Figure 2: VCG transfer payments by agent

```

set.seed(42)
n_sample <- 20

transfer_data <- map_dfr(seq_len(n_sample), function(sim_id) {
  v_build <- runif(n_agents, min = -2, max = 5)
  valuations <- cbind(v_build, rep(0, n_agents))
  vcg_result <- vcg_mechanism(valuations)

  tibble(
    instance = factor(sim_id),
    agent = paste("Agent", seq_len(n_agents)),
    transfer = vcg_result$transfers,
    valuation = v_build
  )
})

p_transfers <- ggplot(transfer_data,
                     aes(x = instance, y = transfer, fill = agent)) +

```




Figure 1: Social welfare under VCG and majority voting compared to the first-best optimum. VCG always achieves the efficient outcome; majority voting sometimes selects the wrong alternative.

```
geom_col(position = position_dodge(width = 0.7), width = 0.6, alpha = 0.85) +
scale_fill_manual(values = okabe_ito[1:3], name = NULL) +
scale_y_continuous(name = "Clarke tax payment") +
labs(x = "Auction instance",
     title = "Transfer Payments in the VCG Mechanism",
     subtitle = "Positive tax = agent is pivotal for the decision") +
theme_publication() +
theme(axis.text.x = element_text(size = 7))
```

p_transfers




Figure 2: Clarke tax payments in the VCG mechanism for a representative sample of 20 instances. Agents who are pivotal — those whose presence changes the chosen outcome — pay a positive tax.

```
save_pub_fig(p_transfers, "vcg-transfers", width = 7, height = 5)
```

Worked example

Three citizens have private valuations for building a public park. The alternative is the status quo (no park, value 0 for all). We apply the VCG mechanism and compare it to majority voting.

```
# Agent valuations for: (Build park, No park)
# Agent 1 loves parks, Agent 2 is mildly positive, Agent 3 prefers no park
valuations <- matrix(c(
  8, 0, # Agent 1: values park at 8
  3, 0, # Agent 2: values park at 3
  -2, 0 # Agent 3: dislikes park (cost = 2)
```

```
), nrow = 3, byrow = TRUE,
  dimnames = list(paste("Agent", 1:3), c("Build", "No build")))

cat("Agent valuations:\n")
```

```
#> Agent valuations:
```

```
print(valuations)
```

```
#>      Build No build
#> Agent 1      8      0
#> Agent 2      3      0
#> Agent 3     -2      0
```

```
# VCG mechanism
vcg <- vcg_mechanism(valuations)
cat(sprintf("\n--- VCG Mechanism ---\n"))
```

```
#>
#> --- VCG Mechanism ---
```

```
cat(sprintf("Chosen outcome: %s\n",
  colnames(valuations)[vcg$chosen_outcome]))
```

```
#> Chosen outcome: Build
```

```
cat(sprintf("Social welfare: %.1f\n", vcg$social_welfare))
```

```
#> Social welfare: 9.0
```

```
cat(sprintf("\nClarke taxes:\n"))
```

```
#>
#> Clarke taxes:

for (i in 1:3) {
  cat(sprintf(" Agent %d: tax = %.1f, net payoff = %.1f\n",
    i, vcg$transfers[i], vcg$payoffs[i]))
}
```

```
#> Agent 1: tax = 0.0, net payoff = 8.0
#> Agent 2: tax = 0.0, net payoff = 3.0
#> Agent 3: tax = 0.0, net payoff = -2.0
```

```
# Majority voting
mv <- majority_voting(valuations)
cat(sprintf("\n--- Majority Voting ---\n"))
```

```
#>
#> --- Majority Voting ---
```

```
cat(sprintf("Chosen outcome: %s\n",
           colnames(valuations)[mv$chosen_outcome]))
```

```
#> Chosen outcome: Build
```

```
cat(sprintf("Social welfare: %.1f\n", mv$social_welfare))
```

```
#> Social welfare: 9.0
```

```
# Verify incentive compatibility: what if Agent 3 lies?
```

```
cat("\n--- Incentive Compatibility Check ---\n")
```

```
#>
```

```
#> --- Incentive Compatibility Check ---
```

```
cat("What if Agent 3 reports valuation -10 instead of -2?\n")
```

```
#> What if Agent 3 reports valuation -10 instead of -2?
```

```
valuations_lie <- valuations
valuations_lie[3, 1] <- -10
vcg_lie <- vcg_mechanism(valuations_lie)
cat(sprintf("Outcome with lie: %s\n",
           colnames(valuations)[vcg_lie$chosen_outcome]))
```

```
#> Outcome with lie: Build
```

```
cat(sprintf("Agent 3's true payoff under lie: %.1f\n",
           valuations[3, vcg_lie$chosen_outcome] - vcg_lie$transfers[3]))
```

```
#> Agent 3's true payoff under lie: -2.0
```

```
cat(sprintf("Agent 3's true payoff under truth: %.1f\n", vcg$payoffs[3]))
```

```
#> Agent 3's true payoff under truth: -2.0
```

```
cat("Lying is not profitable -- VCG is incentive compatible.\n")
```

```
#> Lying is not profitable -- VCG is incentive compatible.
```

The VCG mechanism chooses to build the park (total value $8 + 3 - 2 = 9 > 0$). Agent 3, who opposes the park, pays no Clarke tax because removing Agent 3 would not change the outcome — the other agents still prefer building. Majority voting also selects “Build” (2-1 vote), but this is coincidental; majority voting can fail when a minority has intense preferences that outweigh the mild preferences of the majority.

Extensions

- **Myerson’s optimal auction** (@ref(sec-auctions)) derives the revenue-maximizing mechanism by applying mechanism design to auction settings, introducing virtual valuations and reserve prices.
- **Budget balance** — the VCG mechanism generally runs a deficit (total taxes do not cover costs). The AGV (Arrow–d’Aspremont–Gerard–Varet) mechanism achieves budget balance but only satisfies IC in a Bayesian sense.
- **Combinatorial auctions** extend VCG to settings with multiple items and complementarities, though computational complexity becomes a major concern.
- For the foundational treatment, see @nisan2007 (Chapters 9–12) and @roth2002 on practical market design.

Exercises

1. **Pivotal agents.** In the worked example, modify Agent 2’s valuation for the park to 1 (so total value of building becomes $8 + 1 - 2 = 7$). Recompute the VCG outcome and Clarke taxes. Now set Agent 1’s valuation to 1 so building yields $1 + 1 - 2 = 0$. What happens? Which agents are pivotal, and how do their taxes change?
2. **Budget balance.** Run 500 simulations of the VCG mechanism with 5 agents and a binary public goods decision ($v_i \sim \text{Uniform}(-2, 5)$). Compute the total Clarke tax revenue in each simulation. What fraction of simulations have positive total tax revenue? Plot the distribution of total taxes and explain why VCG generally does not achieve budget balance.
3. **Manipulation under majority voting.** Construct a 3-agent, 3-outcome example where majority voting leads to a Condorcet cycle (no Condorcet winner), while VCG selects a unique efficient outcome. Implement both mechanisms and show that the VCG outcome maximizes social welfare.

Solutions appear in @ref(sec-solutions).