

Probability Essentials

Probability distributions, expectations, Bayes' rule, and stochastic processes.

This appendix reviews the probability concepts that underpin mixed strategies, Bayesian games, and the simulation chapters. The emphasis is on intuition and R implementation rather than measure-theoretic rigour.

Sample spaces and events

A **sample space** Ω is the set of all possible outcomes of a random experiment. An **event** $E \subseteq \Omega$ is a subset of outcomes. A **probability measure** P assigns a number $P(E) \in [0, 1]$ to every event, with $P(\Omega) = 1$.

In game theory, the sample space often corresponds to the set of action profiles or the set of player types.

Discrete distributions

A **discrete random variable** X takes values in a countable set $\{x_1, x_2, \dots\}$ with probability mass function $p(x_i) = P(X = x_i)$.

Common discrete distributions

```
# Bernoulli: coin flip (used in mixed strategies)  
rbinom(10, size = 1, prob = 0.5)
```

```
#> [1] 1 1 0 1 1 1 1 0 1 1
```

```
# Binomial: number of cooperators in n rounds  
dbinom(3, size = 10, prob = 0.6) # P(X = 3)
```

```
#> [1] 0.04246733
```

```
# Uniform: equally likely strategies  
sample(c("Rock", "Paper", "Scissors"), size = 1)
```

```
#> [1] "Scissors"
```

```
# Poisson: arrival counts (used in some mechanism design models)  
rpois(10, lambda = 3)
```

```
#> [1] 4 6 2 3 6 7 1 3 3 5
```

Probability mass function and CDF

```
# PMF and CDF of Binomial(10, 0.4)
x <- 0:10
pmf <- dbinom(x, size = 10, prob = 0.4)
cdf <- pbinom(x, size = 10, prob = 0.4)

tibble(x = x, pmf = pmf, cdf = cdf) |>
  head(6)
```

```
#> # A tibble: 6 x 3
#>       x      pmf      cdf
#>   <int>   <dbl>   <dbl>
#> 1     0 0.00605 0.00605
#> 2     1 0.0403  0.0464
#> 3     2 0.121   0.167
#> 4     3 0.215   0.382
#> 5     4 0.251   0.633
#> 6     5 0.201   0.834
```

Continuous distributions

A **continuous random variable** X has a probability density function $f(x)$ such that $P(a \leq X \leq b) = \int_a^b f(x) dx$.

Common continuous distributions

```
# Uniform on [0, 1]: used for types in Bayesian games
runif(5, min = 0, max = 1)
```

```
#> [1] 0.13871017 0.98889173 0.94666823 0.08243756 0.51421178
```

```
# Normal: payoff noise, error terms
rnorm(5, mean = 0, sd = 1)
```

```
#> [1] -0.2787888 -0.1333213  0.6359504 -0.2842529 -2.6564554
```

```
# Exponential: waiting times, discount factors
rexp(5, rate = 1)
```

```
#> [1] 4.9959685 0.2235573 1.2113841 0.7190924 1.3084920
```

Density, CDF, and quantiles

R uses a consistent naming convention: **d** for density, **p** for CDF, **q** for quantile, and **r** for random draws.

```
# Normal distribution
dnorm(0, mean = 0, sd = 1) # density at 0
```

```
#> [1] 0.3989423
```

```
pnorm(1.96, mean = 0, sd = 1) #  $P(X \leq 1.96)$ 
```

```
#> [1] 0.9750021
```

```
qnorm(0.975, mean = 0, sd = 1) # 97.5th percentile
```

```
#> [1] 1.959964
```

```
rnorm(3, mean = 0, sd = 1) # three random draws
```

```
#> [1] 0.3584021 0.3024309 -0.3941145
```

Expected value and variance

The **expected value** of a discrete random variable is $E[X] = \sum_i x_i p(x_i)$. For a continuous variable, $E[X] = \int x f(x) dx$. In game theory, the expected payoff under a mixed strategy $\mathbf{p}$ is:

$$E[u_i] = \sum_{a \in A} p(a) u_i(a)$$

The **variance** $\text{Var}(X) = E[(X - E[X])^2] = E[X^2] - (E[X])^2$ measures the spread of outcomes.

```
# Expected payoff under a mixed strategy
```

```
payoffs <- c(3, 0, 5, 1)
```

```
probs <- c(0.3, 0.2, 0.1, 0.4)
```

```
# Expected value
```

```
ev <- sum(payoffs * probs)
```

```
ev
```

```
#> [1] 1.8
```

```
# Variance
```

```
variance <- sum(probs * (payoffs - ev)^2)
```

```
variance
```

```
#> [1] 2.36
```

```
# Using simulation to verify
```

```
simulated <- sample(payoffs, size = 100000, replace = TRUE, prob = probs)
```

```
mean(simulated)
```

```
#> [1] 1.80101
```

```
var(simulated)
```

```
#> [1] 2.372757
```

Linearity of expectation

A key property: $E[aX + bY] = aE[X] + bE[Y]$, regardless of whether X and Y are independent. This is used constantly when computing expected payoffs in multi-player games.

Conditional probability

The **conditional probability** of A given B is:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}, \quad P(B) > 0$$

Two events are **independent** if $P(A \cap B) = P(A)P(B)$, equivalently $P(A | B) = P(A)$.

```
# Simulation: P(both cooperate | Player 1 cooperates)
n <- 100000
p1 <- sample(c("C", "D"), n, replace = TRUE, prob = c(0.6, 0.4))
p2 <- sample(c("C", "D"), n, replace = TRUE, prob = c(0.5, 0.5))

# Joint and conditional (assuming independence here)
p_both_C <- mean(p1 == "C" & p2 == "C")
p_p1_C <- mean(p1 == "C")
p_both_C / p_p1_C # should be approx 0.5
```

```
#> [1] 0.5006477
```

Bayes' theorem

Bayes' theorem is the foundation of Bayesian games ([@ref\(sec-bayesian-games\)](#)). It tells us how to update beliefs about a player's type after observing their action:

$$P(\theta | a) = \frac{P(a | \theta) P(\theta)}{P(a)}$$

where θ is the player's type, a is the observed action, $P(\theta)$ is the prior, $P(a | \theta)$ is the likelihood, and $P(a) = \sum_{\theta'} P(a | \theta') P(\theta')$ is the marginal likelihood.

```
# Signaling game: worker chooses Education or No Education
# Types: High (prob 0.6), Low (prob 0.4)
# P(Education | High) = 0.9, P(Education | Low) = 0.2

prior_high <- 0.6
prior_low <- 0.4

p_edu_given_high <- 0.9
p_edu_given_low <- 0.2

# P(Education)
p_edu <- p_edu_given_high * prior_high + p_edu_given_low * prior_low
p_edu
```

```
#> [1] 0.62
```

```
# P(High | Education) by Bayes' theorem
p_high_given_edu <- (p_edu_given_high * prior_high) / p_edu
p_high_given_edu
```

```
#> [1] 0.8709677
```

```
# P(High | No Education)
p_no_edu <- 1 - p_edu
p_high_given_no_edu <- ((1 - p_edu_given_high) * prior_high) / p_no_edu
p_high_given_no_edu
```

```
#> [1] 0.1578947
```

Sequential updating

When multiple signals are observed, Bayes' theorem can be applied sequentially: the posterior from one update becomes the prior for the next. This is used in repeated Bayesian games where players learn about opponents over time.

```
# Sequential Bayesian updating: observe actions round by round
# Prior: P(Cooperator type) = 0.5
# P(C action | Cooperator type) = 0.9
# P(C action | Defector type) = 0.1

update_belief <- function(prior, likelihood_coop, likelihood_defect, action) {
  if (action == "C") {
    l_coop <- likelihood_coop
    l_defect <- likelihood_defect
  } else {
    l_coop <- 1 - likelihood_coop
    l_defect <- 1 - likelihood_defect
  }
  numerator <- l_coop * prior
  denominator <- l_coop * prior + l_defect * (1 - prior)
  numerator / denominator
}

prior <- 0.5
actions <- c("C", "C", "D", "C", "C")

beliefs <- numeric(length(actions))
for (i in seq_along(actions)) {
  prior <- update_belief(prior, 0.9, 0.1, actions[i])
  beliefs[i] <- prior
}

tibble(round = seq_along(actions), action = actions, belief = beliefs)
```

```
#> # A tibble: 5 x 3
```

```
#>   round action belief
#>   <int> <chr>   <dbl>
#> 1     1 C      0.9
#> 2     2 C      0.988
#> 3     3 D      0.9
#> 4     4 C      0.988
#> 5     5 C      0.999
```

Law of large numbers and Monte Carlo

The **law of large numbers** guarantees that the sample mean converges to the true mean as the sample size grows. This justifies the Monte Carlo methods used extensively in [@ref\(sec-monte-carlo\)](#).

```
# Demonstrate convergence of sample mean
n_values <- c(10, 100, 1000, 10000, 100000)
true_mean <- 3.5 # E[Uniform die roll]

estimates <- vapply(n_values, function(n) {
  mean(sample(1:6, n, replace = TRUE))
}, numeric(1))

tibble(n = n_values, estimate = estimates, true_mean = true_mean)
```

```
#> # A tibble: 5 x 3
#>       n estimate true_mean
#>   <dbl>   <dbl>   <dbl>
#> 1    10     3.2     3.5
#> 2   100     3.58    3.5
#> 3  1000     3.48    3.5
#> 4 10000     3.49    3.5
#> 5 100000     3.50    3.5
```

Monte Carlo estimation

To estimate a quantity $\theta = E[g(X)]$, draw N independent samples $X_1, \dots, X_N$ and compute:

$$\hat{\theta}_N = \frac{1}{N} \sum_{i=1}^N g(X_i)$$

The standard error of this estimate is $SE = \sigma/\sqrt{N}$, where σ is the standard deviation of $g(X)$.

```
# Estimate P(at least one cooperator in 5 players) via Monte Carlo
# Each player cooperates independently with probability 0.3

estimate_prob <- function(n_sims, n_players = 5, p_coop = 0.3) {
  counts <- replicate(n_sims, {
    actions <- rbinom(n_players, size = 1, prob = p_coop)
    as.integer(sum(actions) >= 1)
  })
  c(estimate = mean(counts), se = sd(counts) / sqrt(n_sims))
}
```

```

# Exact answer:  $1 - (1 - 0.3)^5$ 
exact <- 1 - 0.7^5

set.seed(42)
mc <- estimate_prob(100000)
tibble(
  method = c("Monte Carlo", "Exact"),
  estimate = c(mc["estimate"], exact)
)

```

```

#> # A tibble: 2 x 2
#>   method      estimate
#>   <chr>         <dbl>
#> 1 Monte Carlo    0.834
#> 2 Exact          0.832

```

Random sampling for games

Several chapters use random sampling to simulate strategic interactions. The key functions are:

```

# sample(): draw from a discrete set
sample(c("Hawk", "Dove"), size = 10, replace = TRUE, prob = c(0.4, 0.6))

```

```

#> [1] "Dove" "Hawk" "Dove" "Dove" "Hawk" "Hawk" "Dove" "Hawk" "Dove" "Hawk"

```

```

# replicate(): repeat a simulation
replicate(5, sum(sample(1:6, 2, replace = TRUE)))

```

```

#> [1] 4 6 5 6 7

```

```

# set.seed(): ensure reproducibility
set.seed(123)
runif(3)

```

```

#> [1] 0.2875775 0.7883051 0.4089769

```

```

set.seed(123)
runif(3) # identical output

```

```

#> [1] 0.2875775 0.7883051 0.4089769

```

Summary of R probability functions

Table 1: R functions for common probability distributions

Distribution	PMF/PDF	CDF	Quantile	Random
Binomial	<code>dbinom()</code>	<code>pbinom()</code>	<code>qbinom()</code>	<code>rbinom()</code>
Poisson	<code>dpois()</code>	<code>ppois()</code>	<code>qpois()</code>	<code>rpois()</code>
Uniform	<code>dunif()</code>	<code>punif()</code>	<code>qunif()</code>	<code>runif()</code>
Normal	<code>dnorm()</code>	<code>pnorm()</code>	<code>qnorm()</code>	<code>rnorm()</code>
Exponential	<code>dexp()</code>	<code>pexp()</code>	<code>qexp()</code>	<code>rexp()</code>